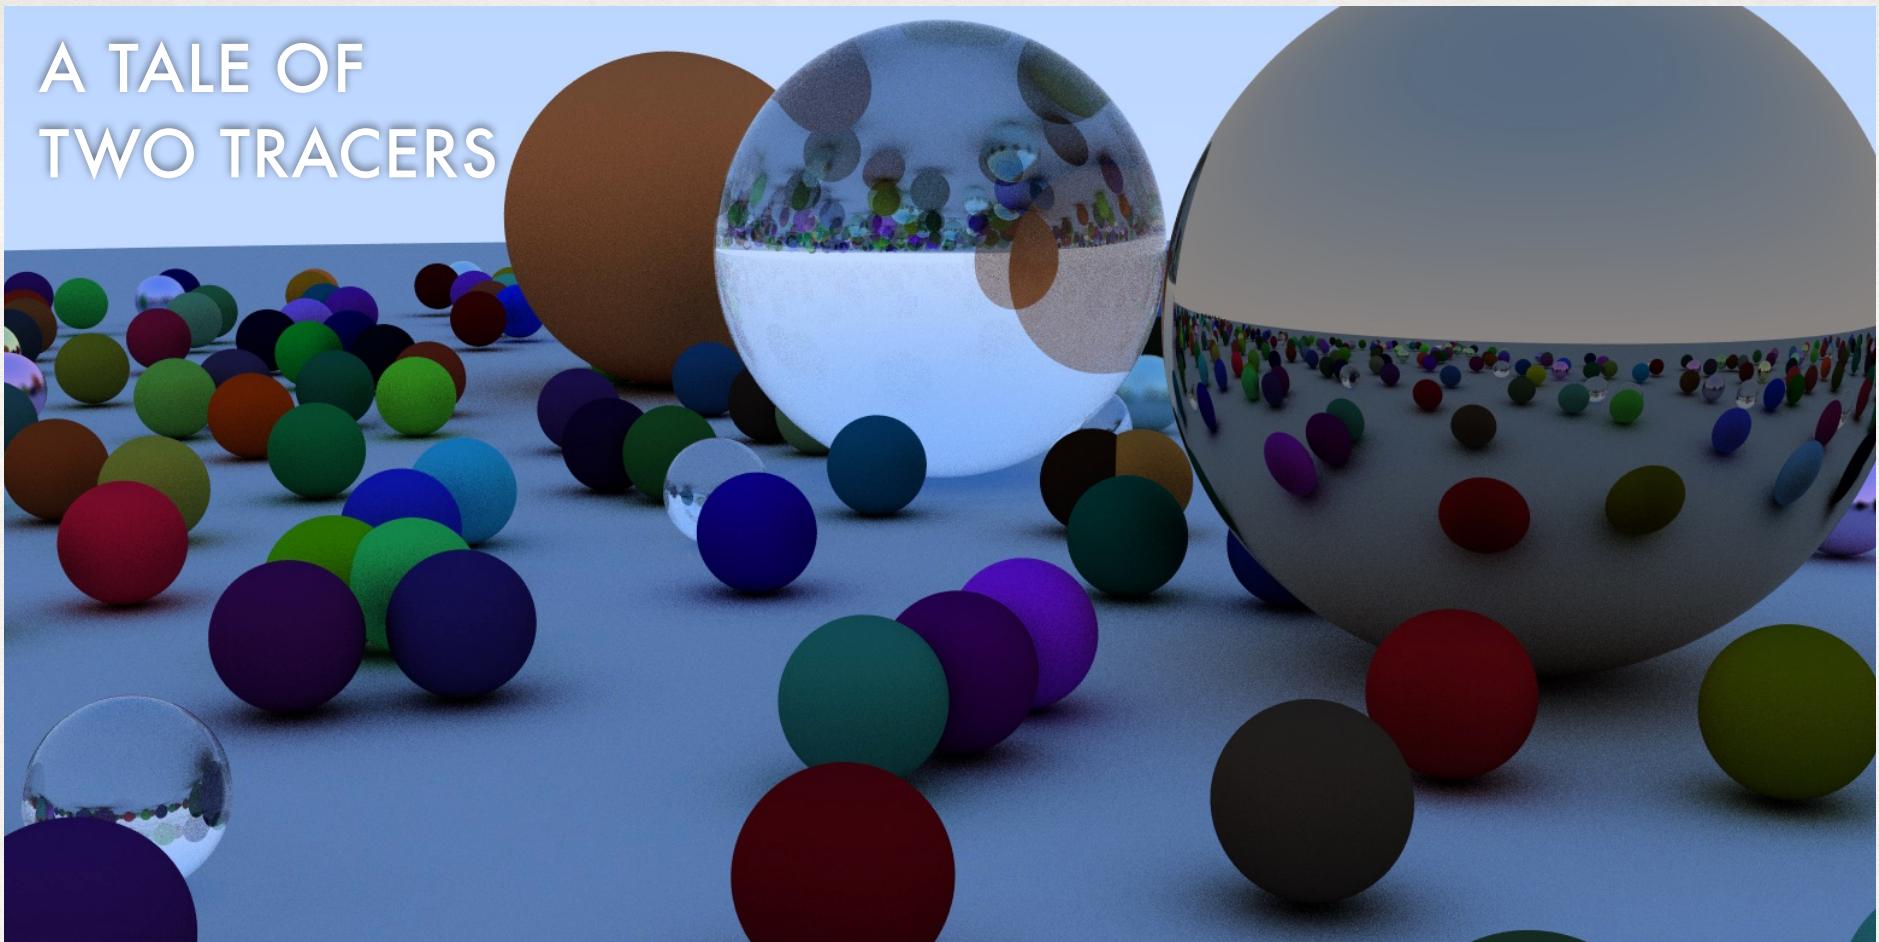


A TALE OF TWO TRACERS



Jonathan Blocksom
@jblocksom
Capital One



Jeff Biggus
@hyperjeff
<http://blog.hyperjeff.net>

360iDev 2016
August 23
Denver, CO

WHY?

IMPLEMENT A RAY TRACER IN SWIFT AND AGAIN IN METAL

- Ray Tracing:

Well defined problem, tons of prior art, resources, simple, extendable

- Why Swift?

It's new (or was when we started)

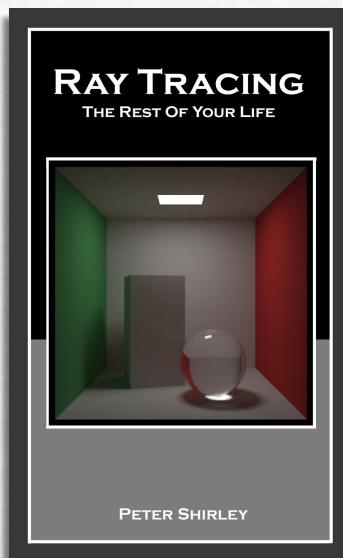
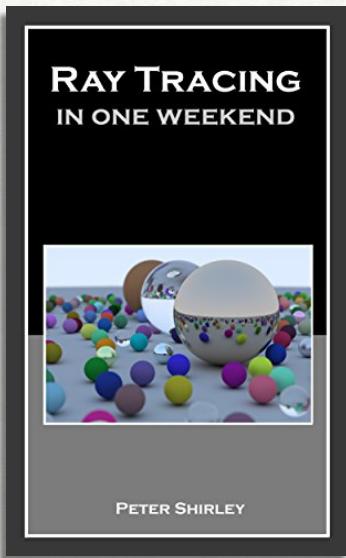
- Why Metal?

It's new and on the GPU

- Why Both?

Compare type-safe, fancy approach with performance oriented GPU approach

QUICK SHOUT OUT

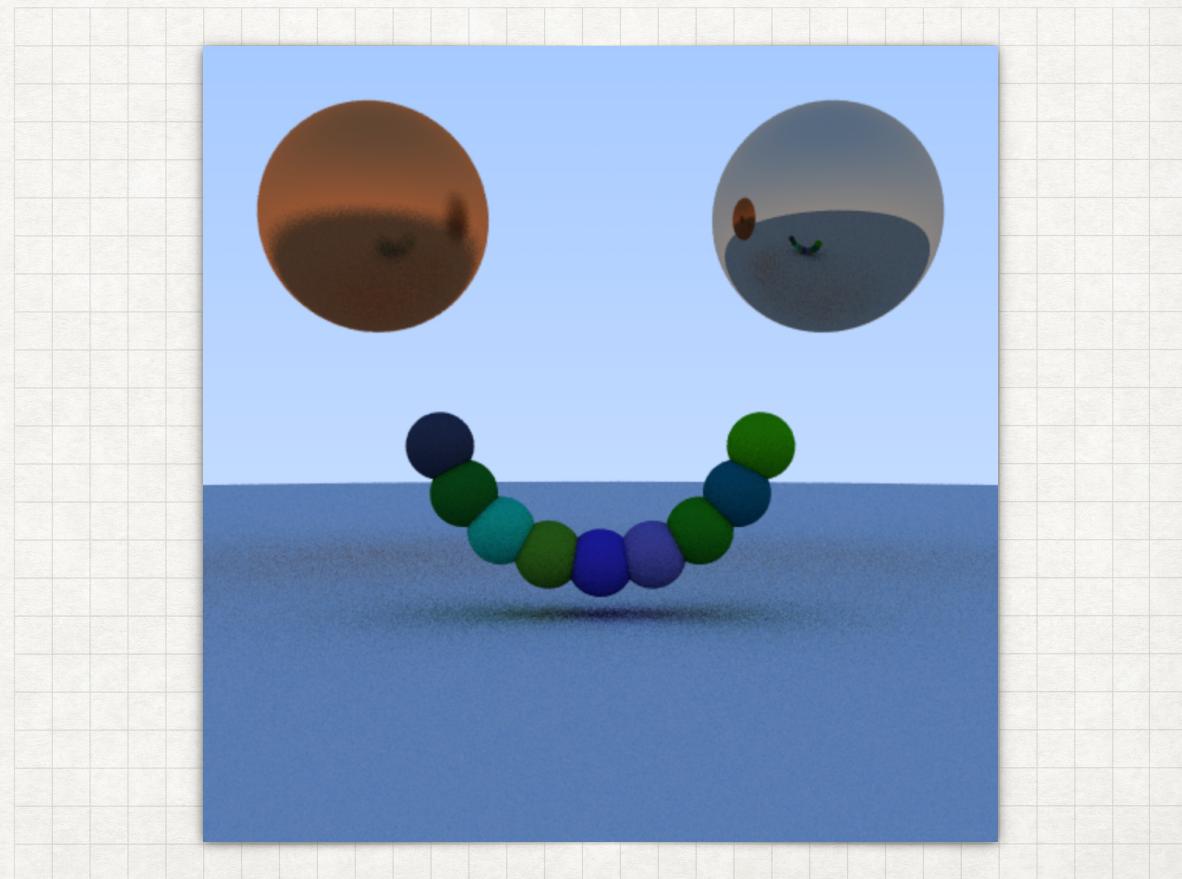


Peter Shirley — @Peter_shirley, <http://psgraphics.blogspot.com>
<https://www.amazon.com/Ray-Tracing-Weekend-Minibooks-Book-ebook/dp/B01B5AODD8>

AGENDA

WE HAVE A PLAN

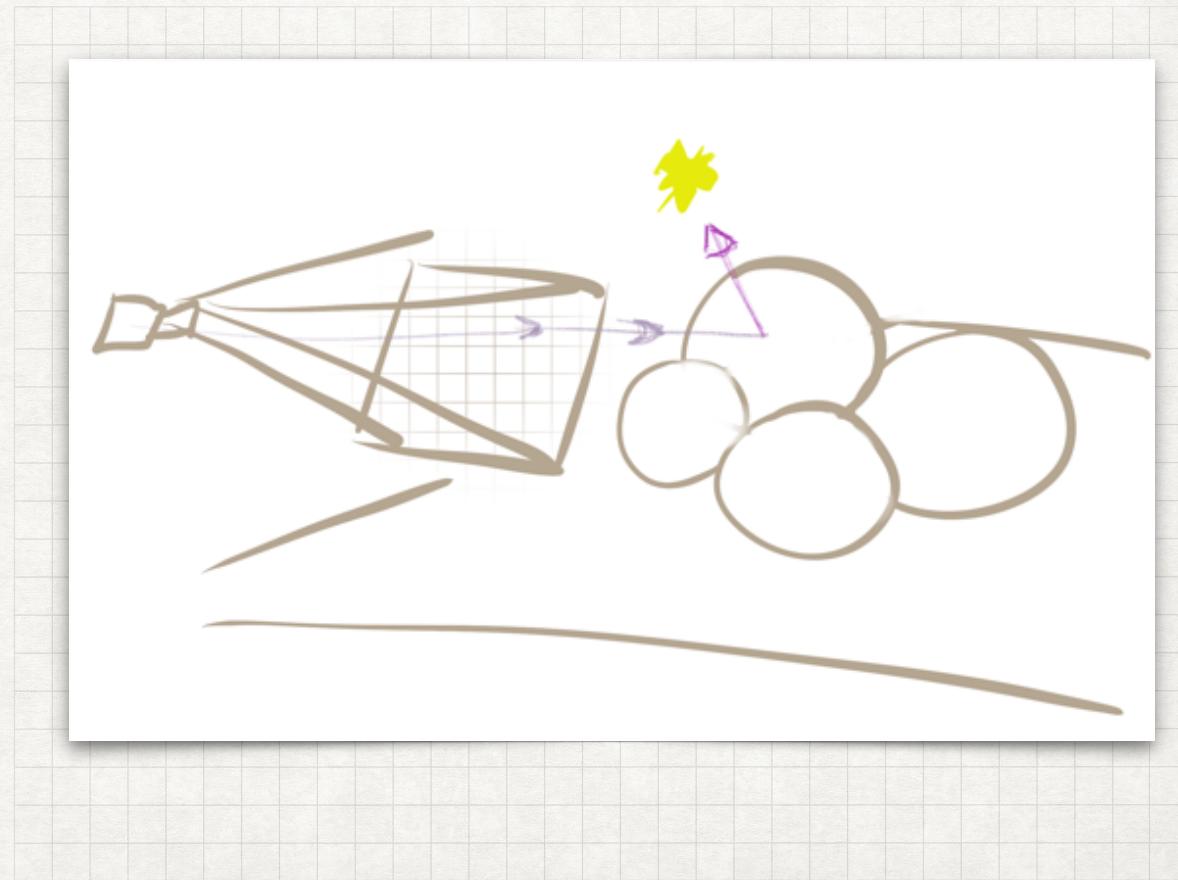
- How Ray Tracing works
- Swift Ray Tracer
- Metal Ray Tracer



RAY TRACING

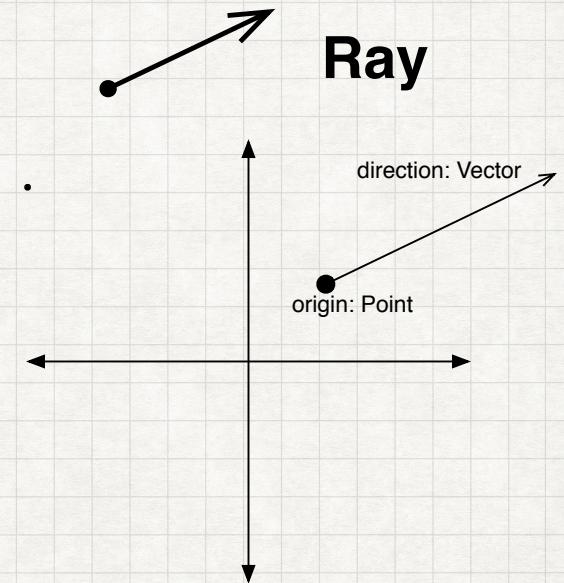
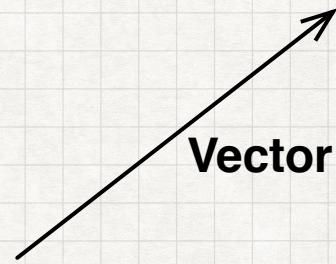
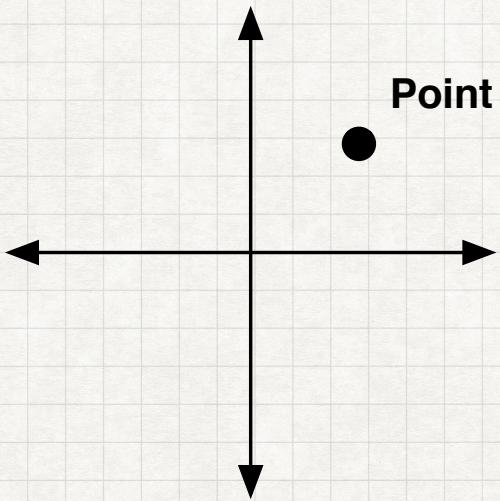
HOW DOES IT WORK?

- Generate an image by creating a ray from the camera to each pixel seeing what it hits
- Determine color of what it hits by shooting new ray from camera ray intersection point
- Ray tracing: work from eye
Conventional 3D: work from objects



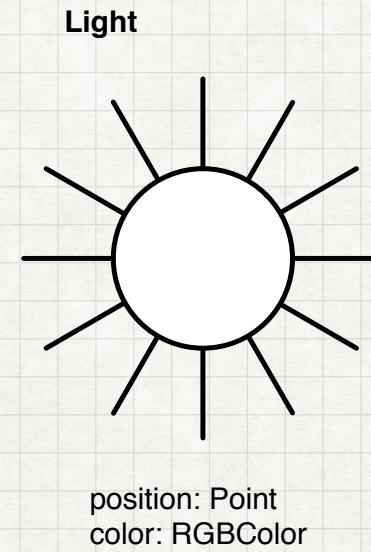
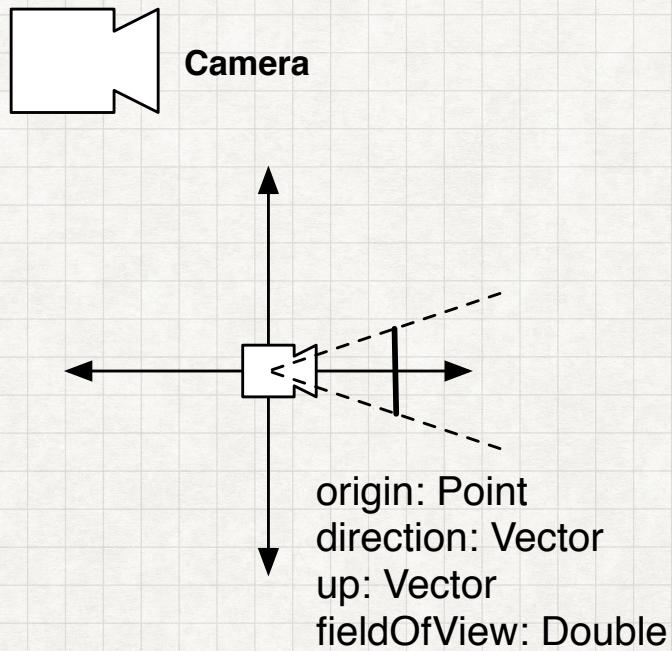
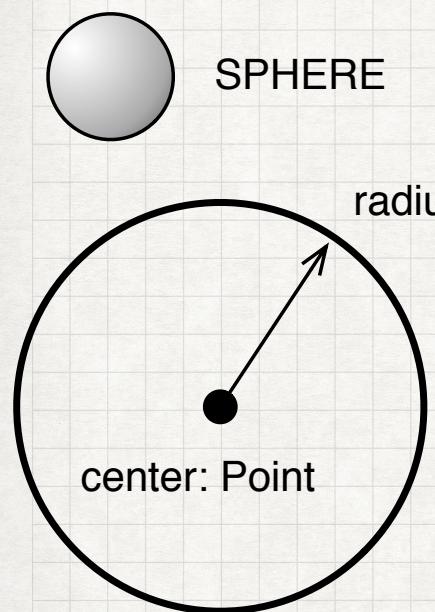
THINGS WE NEED

POINT, VECTOR, RAY



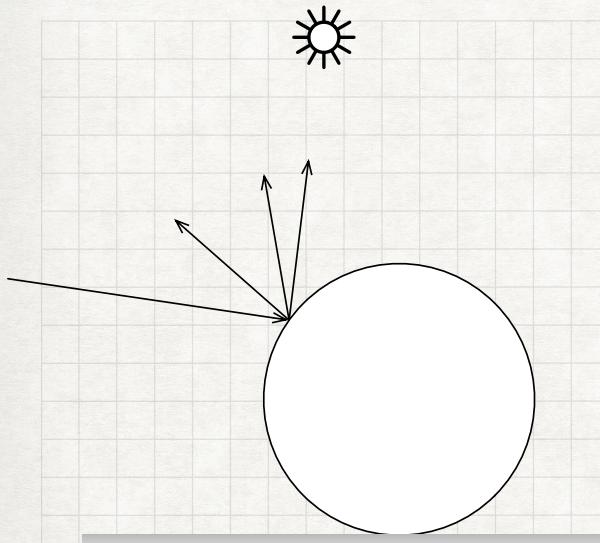
THINGS WE NEED

CAMERA, SPHERE, LIGHT



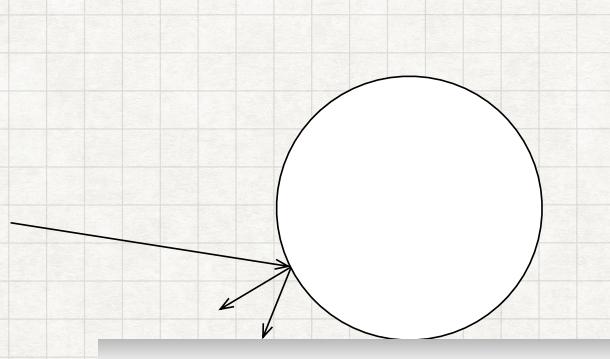
THINGS WE NEED

MATERIALS



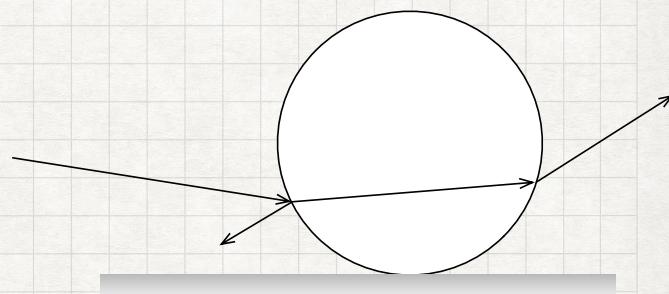
Diffuse

Shade based on color,
direction to light



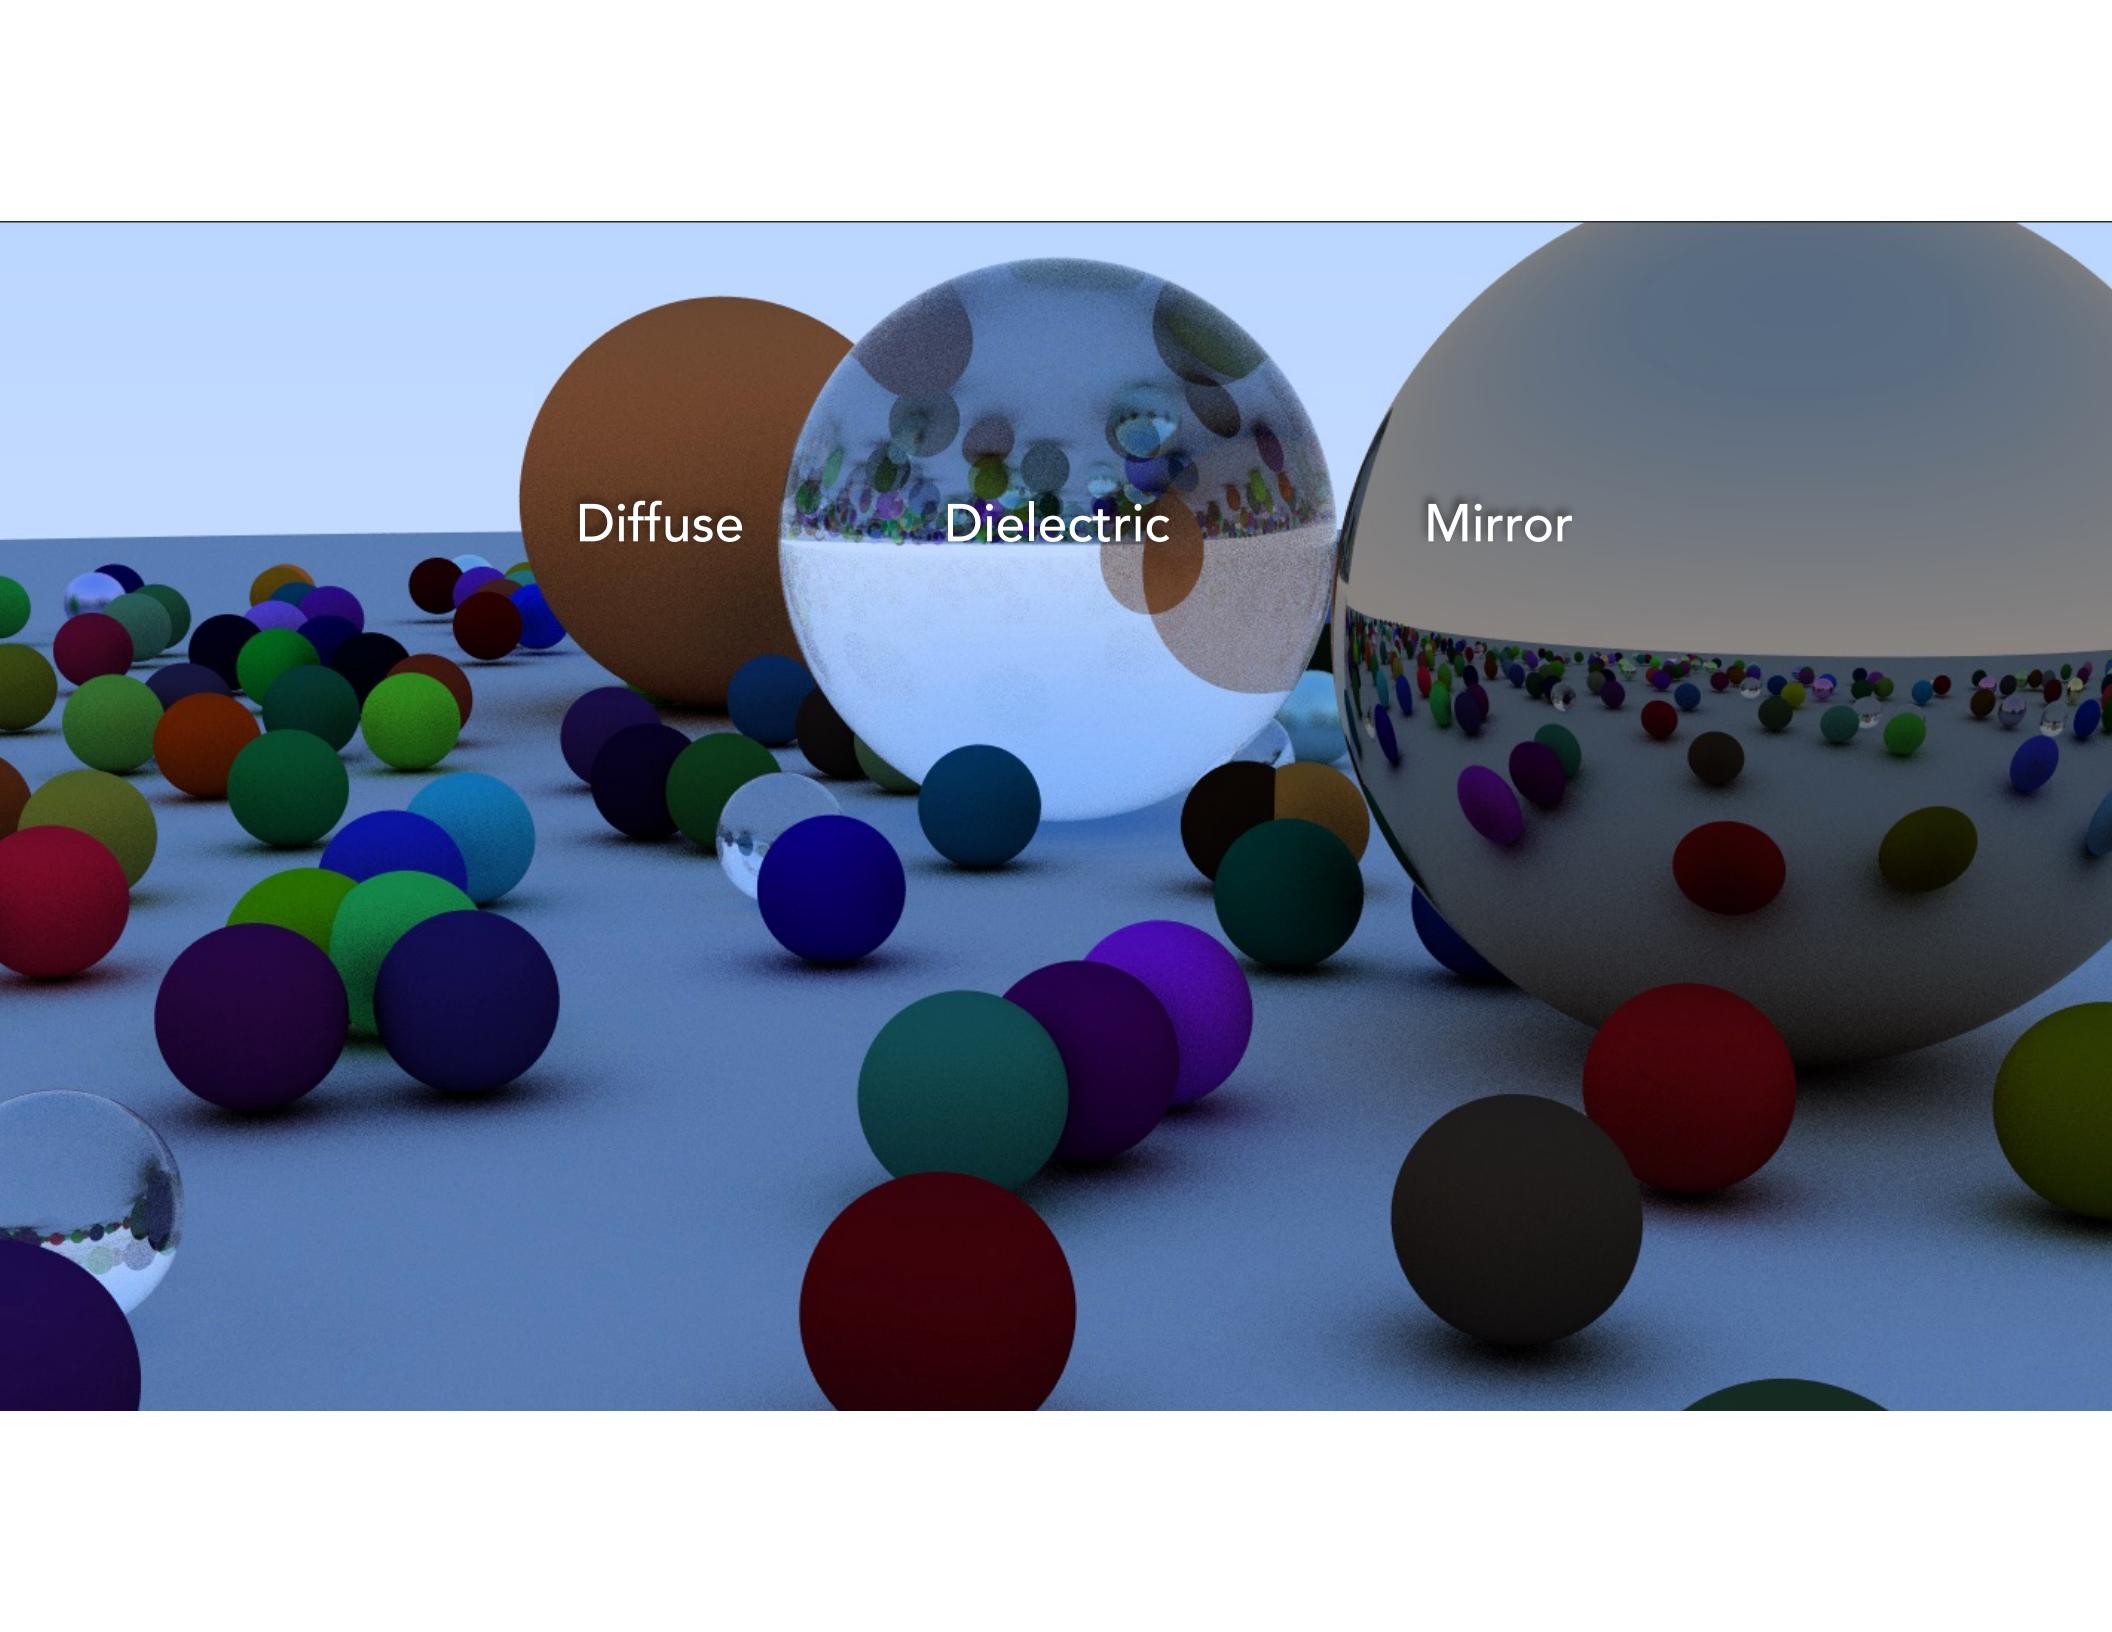
Mirror

Bounce ray, shade based on where
it hits
May attenuate



Dielectric

Ray may go through object



Diffuse

Dielectric

Mirror

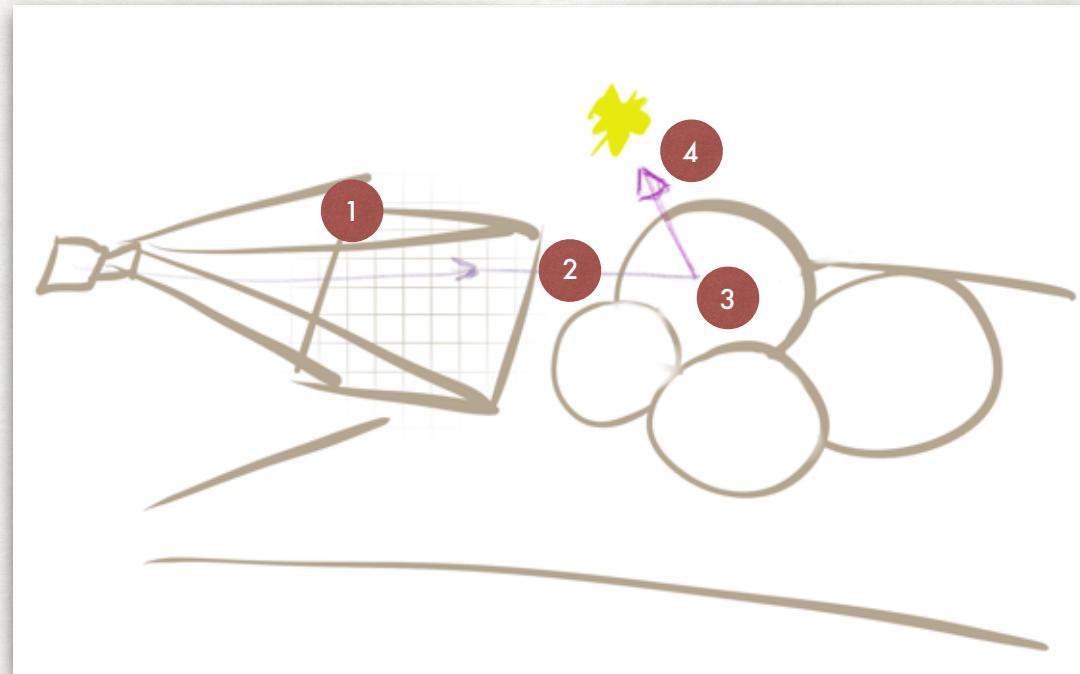
BASIC ALGORITHM

DONEC QUIS NUNC

- 1 For all the pixels in the frame
- 2 Generate a ray through the screen
- 3 Figure out what the ray hits
- 4 Figure out the color of surface

Computational Complexity:

Number of pixels * number of objects

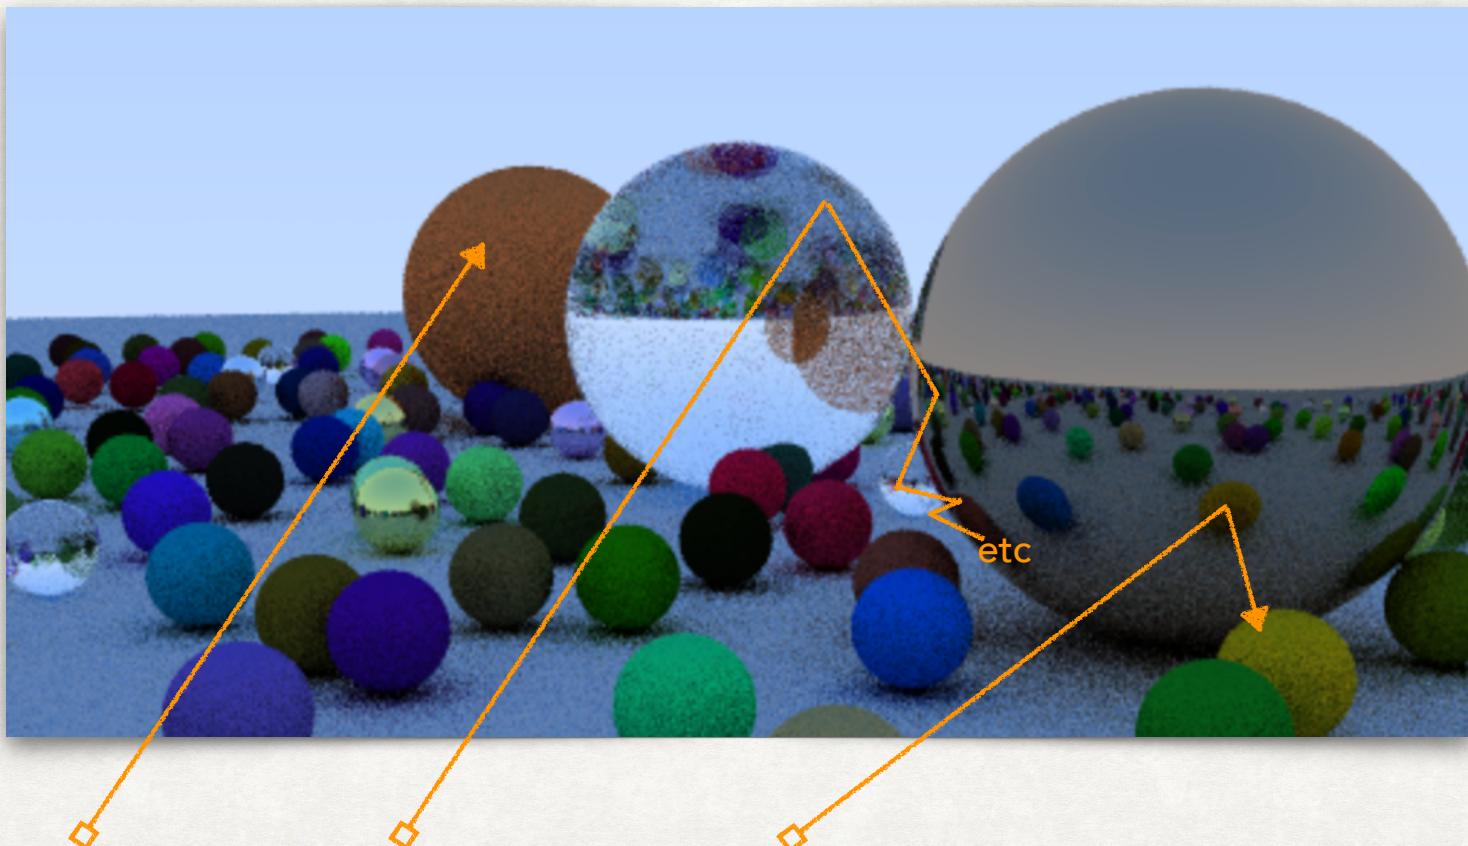


BETTER IMAGES

DONEC QUIS NUNC

- Bounce Depth
- Anti-Aliasing

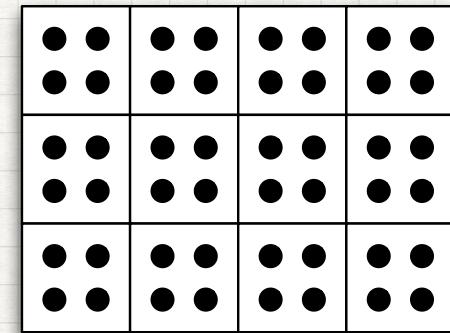
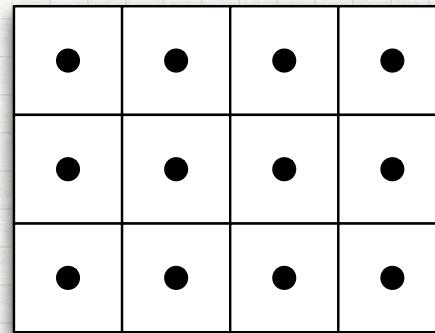
BOUNCE DEPTH



ANTI-ALIASING

SAMPLES PER PIXEL

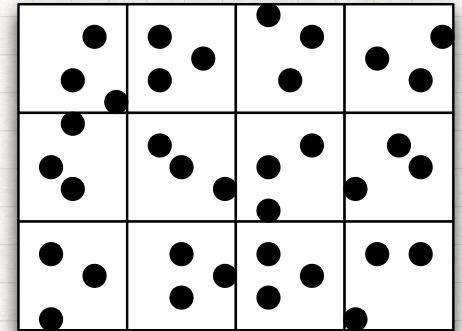
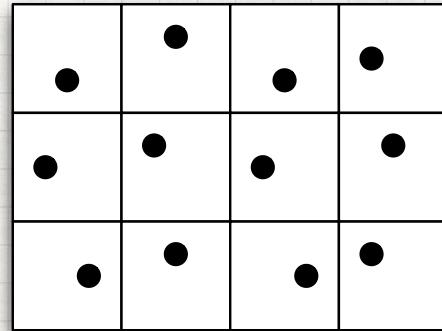
- One sample per pixel causes jaggies
- Multiple samples smooth out edges



ANTI-ALIASING

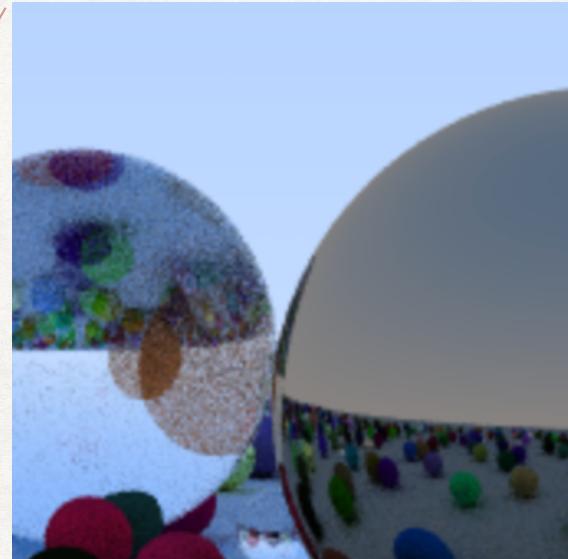
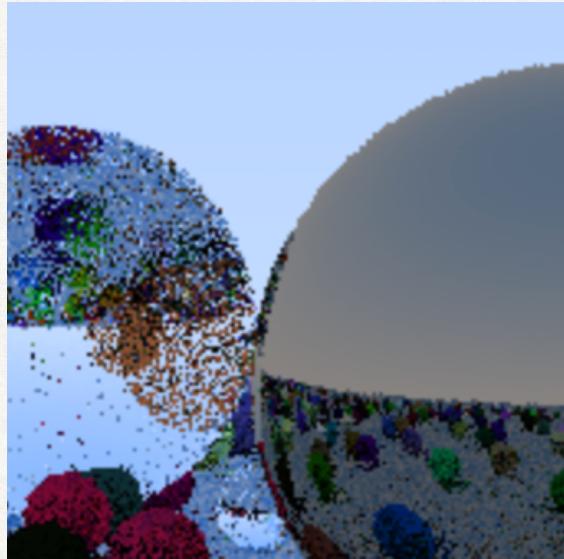
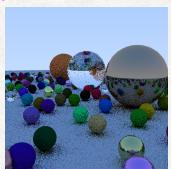
JITTERED SAMPLES

- Avoid regular patterns in objects
- Add random offset to each ray



ANTI-ALIASING RESULTS

- From one sample to 20 samples



A SWIFT RAY TRACER

GOALS

- Try to use value types, protocols
- Functional where appropriate
- Not very concerned with rendering time
- One and a half of them!

GEOMETRIC TOOLBOX

MATH IS FUN

- Point
- Vector
- Rays
- Objects
- Ray / Object Intersection

VECTOR / POINT: ONE OR TWO TYPES?

- Both are just three values: x, y, z
- Can use Swift type safety to make sure your math is right
- Make conversions explicit
- Won't accidentally use ray direction (vector) instead of ray origin (point)

```
public struct Ray {  
    let origin: Point  
    let direction: Vector  
  
    init(origin: Point, direction: Vector) {  
        self.origin = origin  
        self.direction = normalize(direction)  
    }  
  
    func pointAlong(distance: Double) -> Point {  
        return origin + distance * direction  
    }  
}
```

VECTOR AND POINT STRUCTS

TWO TYPES VERSION

```
public struct Vector {
    let x, y, z: Double

    public init(_ x: Double, _ y: Double, _ z: Double) {
        self.x = x
        self.y = y
        self.z = z
    }

    func length () -> Double {
        let magSquared = x*x + y*y + z*z
        let length = sqrt(magSquared)
        return length
    }

    public func normalize(v: Vector) -> Vector {
        let vLen = v.length()
        if vLen > 0.0 {
            return (1.0 / vLen) * v
        } else {
            return v
        }
    }

    public func +(v1: Vector, v2: Vector) -> Vector {
        return Vector(v1.x + v2.x, v1.y + v2.y, v1.z + v2.z)
    }

    public prefix func -(v1: Vector) -> Vector {
        return Vector(-v1.x, -v1.y, -v1.z)
    }

    public func -(v1: Vector, v2: Vector) -> Vector {
        return v1 + (-v2)
    }

    infix operator . { associativity left precedence 150 }
    public func .(v1: Vector, v2: Vector) -> Double {
        return (v1.x * v2.x + v1.y * v2.y + v1.z * v2.z)
    }

    public func *(a: Double, v: Vector) -> Vector {
        return Vector(a * v.x, a * v.y, a * v.z)
    }
}
```

```
public struct Point {
    let x, y, z: Double

    public init(_ x: Double, _ y: Double, _ z: Double) {
        self.x = x
        self.y = y
        self.z = z
    }

    func relativeTo(origin: Point) -> Point {
        return Point(x - origin.x, y - origin.y, z - origin.z)
    }

    func vectorTo(p: Point) -> Vector {
        return Vector(p.x - x, p.y - y, p.z - z)
    }

    func distanceBetween(p: Point) -> Double {
        return self.vectorTo(p).length()
    }

    let Origin = Point(0, 0, 0)

    public func +(p: Point, v: Vector) -> Point {
        return Point(p.x + v.x, p.y + v.y, p.z + v.z)
    }

    // a ~ b  a0 a1 b2*
    infix operator — { associativity left precedence 150 }
    func —(p1: Point, p2: Point) -> Vector {
        return Vector(p2.x - p1.x, p2.y - p1.y, p2.z - p1.z)
    }
}
```

SWIFT + SIMD

DONEC QUIS NUNC

- Expressiveness of Swift,
Speed of SIMD library
- Implements operators, vector operations

```
import simd  
  
public typealias Vector = float3
```

RAY - OBJECT INTERSECTION

- Optionals shine!

```
public override func hitTest(passingRay: Ray, limits: FloatRange) -> HitResponse? {  
    ...  
}
```

- Easy to keep track of intersections

```
public func castRay(r: Ray) -> RaySurfaceIntersection? {  
  
    var closestIntersection: RaySurfaceIntersection? = nil  
    for object in objects {  
        if let currentIntersection = object.intersect(r) {  
            if closestIntersection == nil {  
                closestIntersection = currentIntersection  
            } else {  
                if currentIntersection.closerThan(closestIntersection!) {  
                    closestIntersection = currentIntersection  
                }  
            }  
        }  
    }  
  
    return closestIntersection  
}
```

WITHOUT OPTIONALS

- C++ implementation from Ray Tracing in One Weekend

```
#ifndef HITABLELISTH
#define HITABLELISTH

#include "hitable.h"

class hitable_list: public hitable {
public:
    hitable_list() {}
    hitable_list(hitable **l, int n) {list = l; list_size = n; }
    virtual bool hit(const ray& r, float tmin, float tmax, hit_record& rec) const;
    hitable **list;
    int list_size;
};

bool hitable_list::hit(const ray& r, float t_min, float t_max, hit_record& rec) const {
    hit_record temp_rec;
    bool hit_anything = false;
    double closest_so_far = t_max;
    for (int i = 0; i < list_size; i++) {
        if (list[i]->hit(r, t_min, closest_so_far, temp_rec)) {
            hit_anything = true;
            closest_so_far = temp_rec.t;
            rec = temp_rec;
        }
    }
    return hit_anything;
}

#endif
```

RAY-SPHERE INTERSECTION

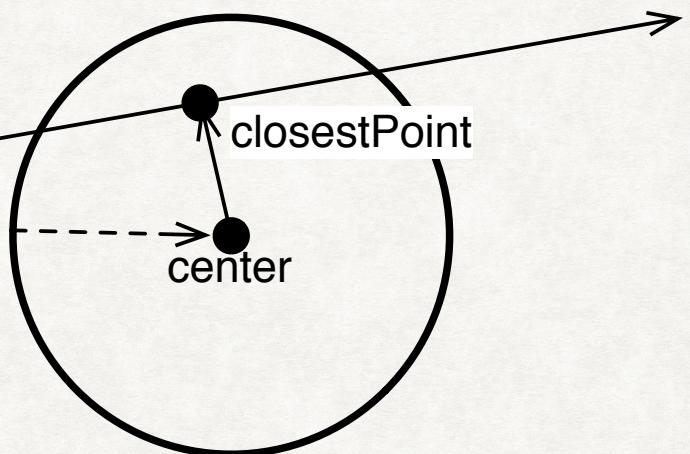
ONE WAY: 1/3

- Project vector to sphere onto ray direction



```
// Get origin relative to sphere center
let rayOriginToSphere = r.origin → center

// Calculate point along ray of closest approach to sphere center
let closestDistance = rayOriginToSphere · r.direction
if (closestDistance < 0) {
    // Sphere is behind the ray, no intersection
    return nil
}
let closestPoint = r.pointAlong(closestDistance)
```

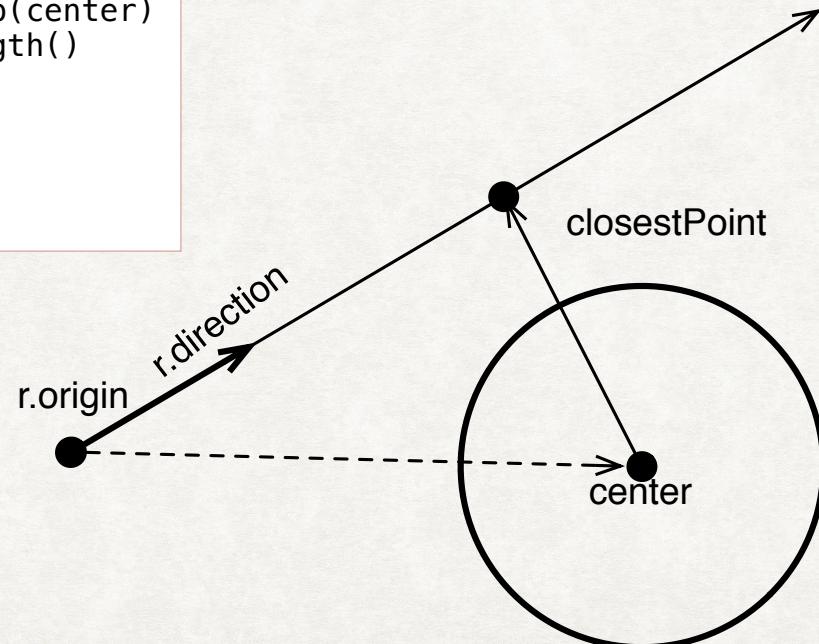


RAY-SPHERE INTERSECTION

ONE WAY: 2/3

- Reject rays outside of sphere

```
// Calculate distance of closest approach to center
let closestPointToCenter = closestPoint.vectorTo(center)
let distanceToCenter = closestPointToCenter.length()
if (distanceToCenter > radius) {
    // Sphere does not hit the ray
    return nil
}
```



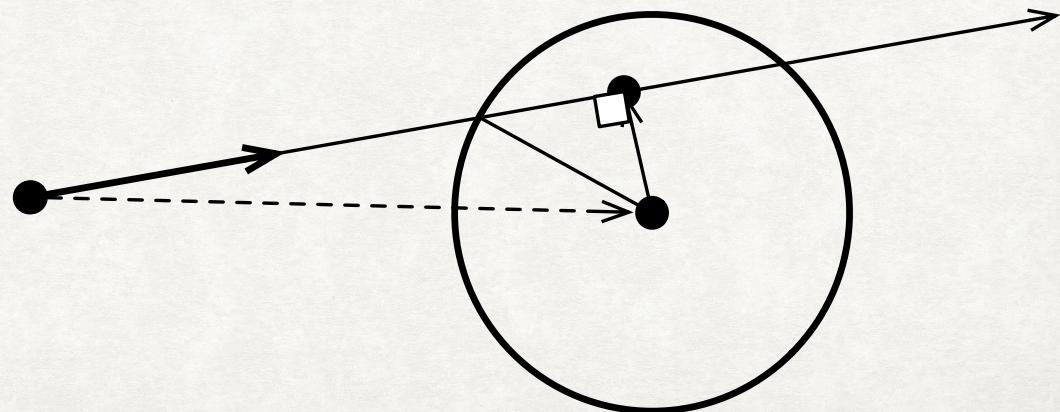
RAY-SPHERE INTERSECTION

ONE WAY: 3/3

- Find point on sphere using Pythagorean Theorem

```
// Ray intersects sphere, figure out where
// Calculate distance between intersection and closest approach
let distanceToIntersection = sqrt(radius*radius - distanceToCenter*distanceToCenter)
let directionSign = (distanceToIntersection > closestDistance) ? -1.0 : 1.0

let intersectionDistance = closestDistance - directionSign * distanceToIntersection
let intersectionPoint = RayPoint(r, intersectionDistance)
```



RAY - SPHERE INTERSECTION

ANOTHER WAY

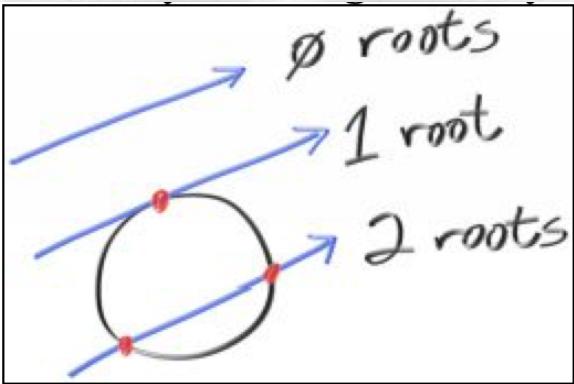
- Equation of ray:

$$p = \text{ray.origin} + t * \text{ray.direction}$$

- Equation of sphere:

$$(x-cx)^2 + (y-cy)^2 + (z-cz)^2 = R^2$$

- Solve simultaneous equations



```
let centerToRayOrigin = passingRay.origin - center
let c = vector_length_squared(centerToRayOrigin) - radius * radius

let rayOriginWithinSphere = 0 < c
let rayDirectionTowardSphere = 0 < dot(passingRay.direction, centerToRayOrigin)

guard rayOriginWithinSphere || rayDirectionTowardSphere else { return nil }

let a = vector_length_squared(passingRay.direction)
let b = dot(centerToRayOrigin, passingRay.direction)

let discriminant = b * b - a * c

guard 0 <= discriminant else { return nil }

let partA = -1 * b / a
let partB = sqrt(discriminant) / a
var solution = partA - partB

if !limits.contains(solution) {
    solution = partA + partB
}

guard limits.contains(solution) else { return nil }

let surfaceIntersection = passingRay.origin + passingRay.direction * solution
```

UNICODE FUN!

LET 😊=TRUE

- $a \rightarrow b$

Vector from two points

- $a \cdot b$

Dot product

- 180°

Degrees to radians

- let $\pi = \text{Float}(M_PI)$

- In variable names

```
let v|| = dot(direction, upVector) * upVector
let v_ = direction - v||
θ = atanf(v_.x/v_.z)

φ = acosf( length(v_) / length(direction) )
vᵀᴹ = normalize(cross(v_, direction))
φRotationRequired = 0.0000001 < abs(φ)
```

ENUMS AND PROTOCOLS

- Hittable

```
public class Hittable {  
  
    public var material: Material = .lambertian  
  
    func hitTest(passingRay: Ray, limits: FloatRange) -> HitResponse? {  
        return nil  
    }  
}
```

- HitResponse

```
public enum HitResponse {  
    case bounce(distance: Float, ray: Ray, colorReflectivity: Color)  
    case color(distance: Float, value: Color)  
}
```

- Material

```
public enum Material {  
    case lambertian  
    case metal(fuzz: Float)  
    case dialectric(indexOfRefraction: Float)  
  
    func lightResponse(passingRay: Ray, normal: Vector,  
                      surfaceIntersection: Vector, colorReflectivity: Color,  
                      randomInteriorPoint: Vector) -> HitResponse?
```

RENDERING LOOP DESIGN

- Cast rays
- Light results

```
func lightResponse(...) -> HitResponse? {
    let length = distance(passingRay.origin, surfaceIntersection)
    let reflected = vector_reflect(passingRay.direction, normal)

    switch self {

        case .lambertian:
            let nextDirection = normal + randomInteriorPoint

            return .bounce(distance: length,
                           ray: Ray(origin: surfaceIntersection, direction: nextDirection),
                           colorReflectivity: colorReflectivity)

        case let .metal(fuzz):
            let nextDirection = reflected + min(fuzz, 1) * randomInteriorPoint

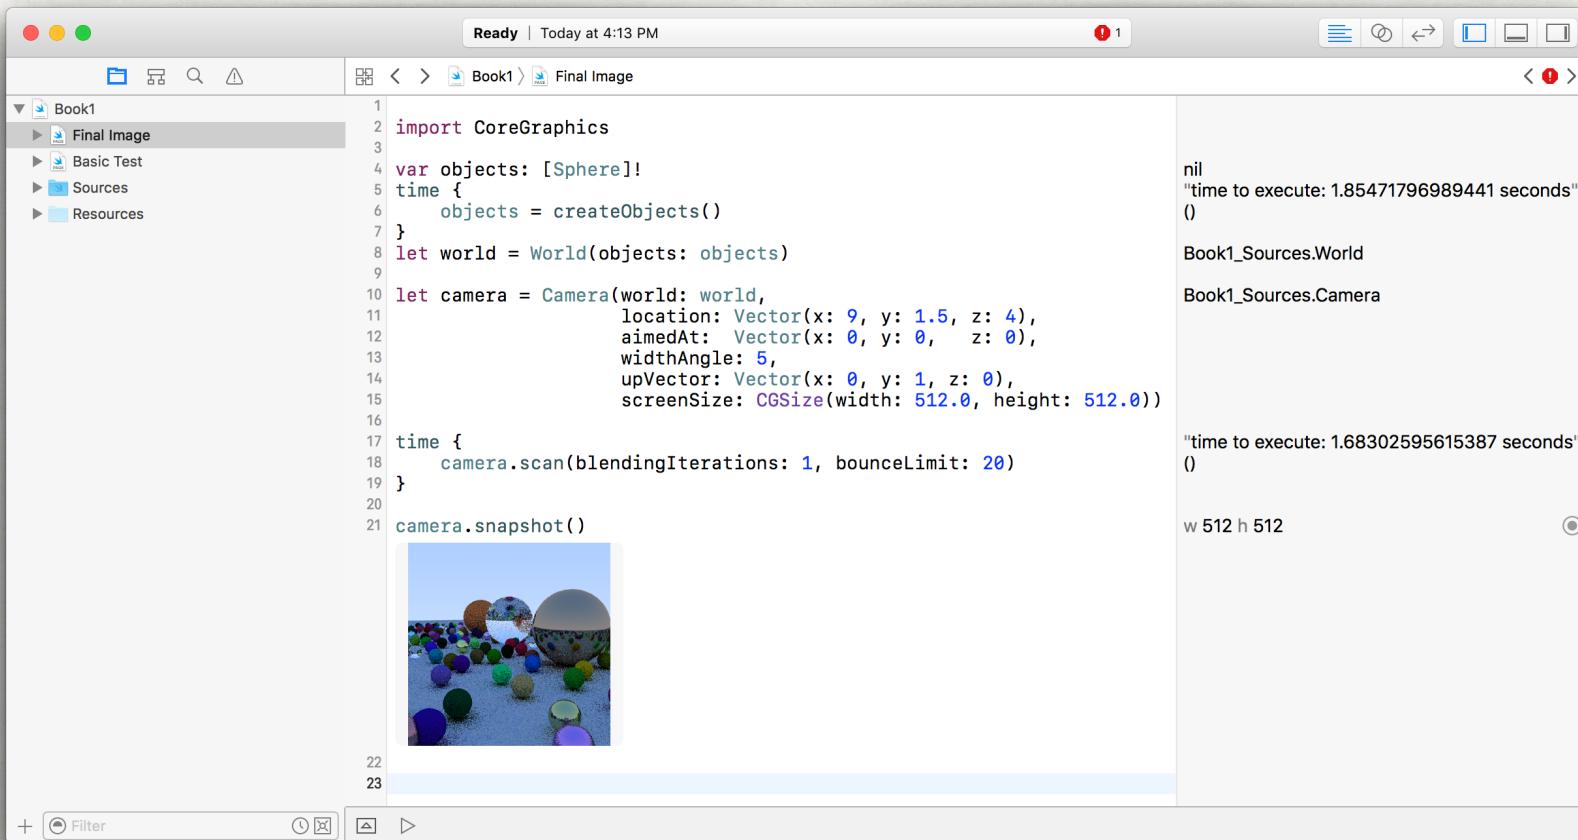
            guard 0 < dot(nextDirection, normal) else { return nil }

            return .bounce(distance: length,
                           ray: Ray(origin: surfaceIntersection, direction: nextDirection),
                           colorReflectivity: colorReflectivity)

        case let .dielectric(indexOfRefraction):
            let attenuation = Color.white()
            ...

            return .bounce(distance: length,
                           ray: Ray(origin: surfaceIntersection, direction: reflected),
                           colorReflectivity: attenuation)
    }
}
```

ALL IN A PLAYGROUND



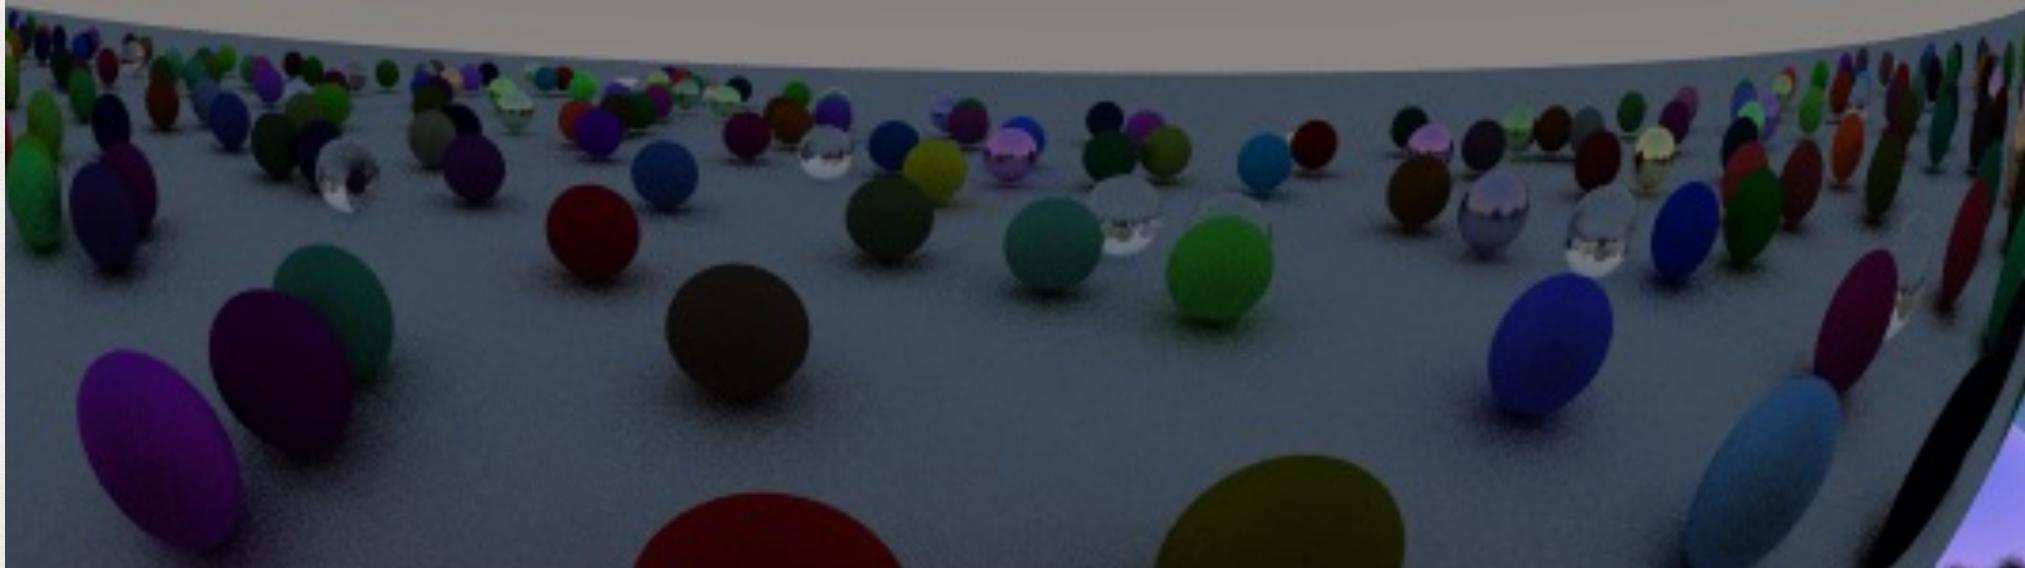
A screenshot of a Mac OS X application window titled "Ready | Today at 4:13 PM". The window contains a code editor and a preview pane. The code editor shows a Swift script named "Final Image" under the "Book1" project. The script imports CoreGraphics and defines a "World" object containing a "Camera" and a "snapshot" method. The preview pane shows a 3D rendering of several colored spheres (red, green, blue, yellow) on a reflective surface.

```
1 import CoreGraphics
2
3 var objects: [Sphere]!
4 time {
5     objects = createObjects()
6 }
7 let world = World(objects: objects)
8
9 let camera = Camera(world: world,
10                     location: Vector(x: 9, y: 1.5, z: 4),
11                     aimedAt: Vector(x: 0, y: 0, z: 0),
12                     widthAngle: 5,
13                     upVector: Vector(x: 0, y: 1, z: 0),
14                     screenSize: CGSize(width: 512.0, height: 512.0))
15
16 time {
17     camera.scan(blendingIterations: 1, bounceLimit: 20)
18 }
19
20 camera.snapshot()
```

The preview pane displays the resulting image, which is a 3D rendering of several spheres of different colors (red, green, blue, yellow) scattered on a reflective surface. The output of the code is also shown in the preview pane:

```
nil
"time to execute: 1.85471796989441 seconds"
()
Book1_Sources.World
Book1_Sources.Camera
"time to execute: 1.68302595615387 seconds"
()
w 512 h 512
```

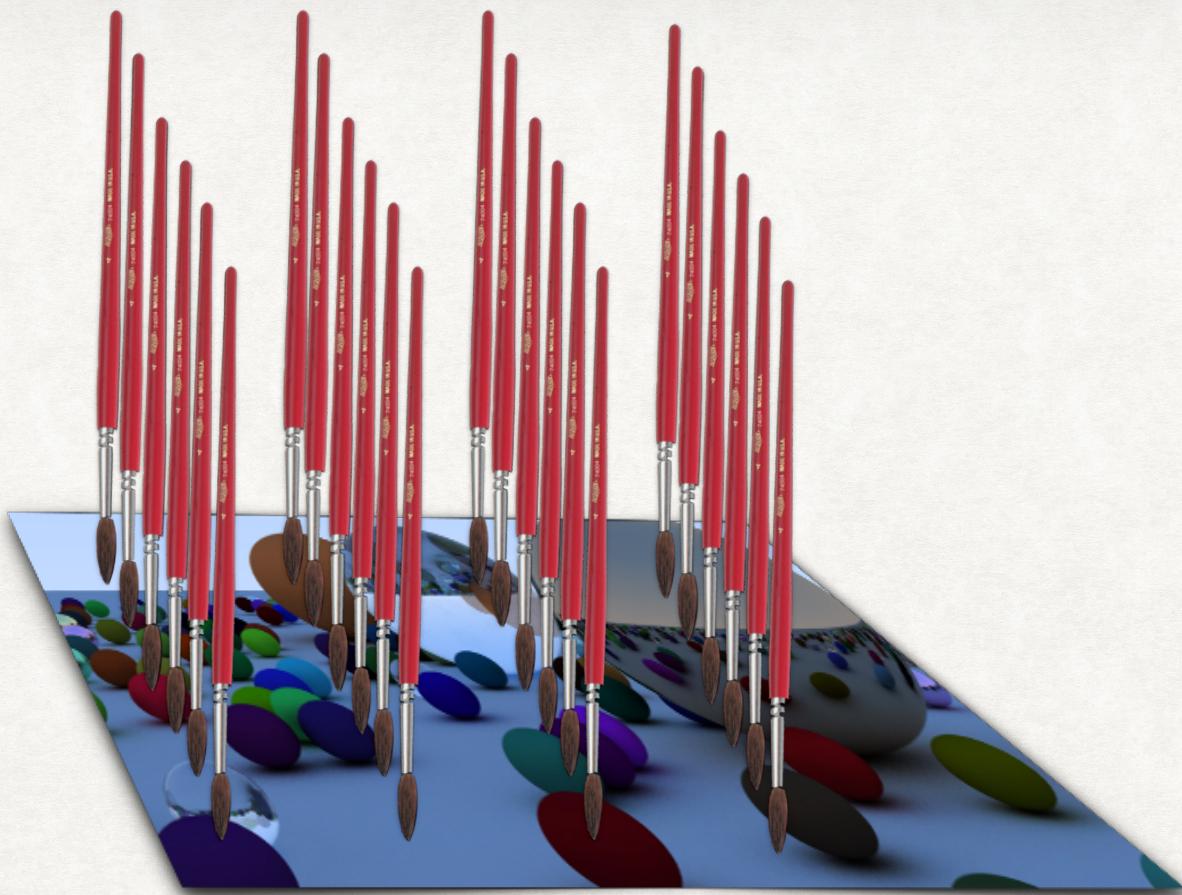
A METAL RAY TRACER



IDEA

- Using GPUs can hold out the hope of big speed gains
- (Perhaps hardest part) What is parallel about your problem?
- Set up shop in Metal to do each individual calculation
- Metal is a variant of C++-14, with much of the evil stuff stripped out
- Prepare what all Metal threads need in common, and pass that from the host
- Figure out how you need to get the results back

```
kernel void rayShader( coordinate ) {  
    calculateColorForRayAtCoordinate(x,y);  
    return color;  
}
```



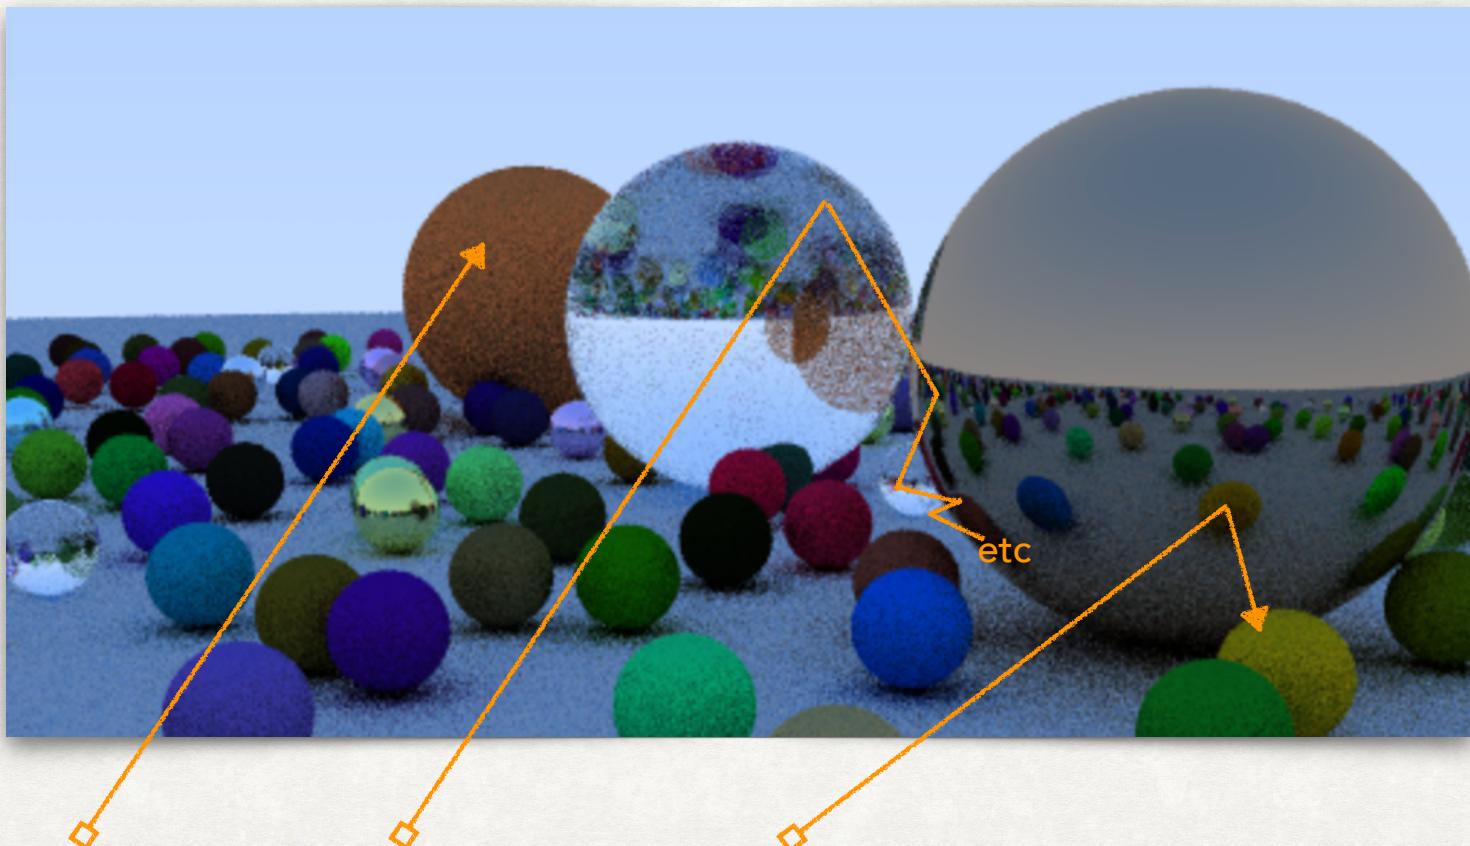
```
kernel void rayShader(const device Sphere* objectsIn [[ buffer(0) ]],  
                      const device ViewInfo* viewInfoIn [[ buffer(1) ]],  
  
                      texture2d<float, access::write> output [[texture(0)]],  
                      COLOR = 6  
                      uint2 upos [[thread_position_in_grid]] ) {  
    PLOT 12, 7  
  
    // stuff!  
  
    output.write(float4(totalColor.x, totalColor.y, totalColor.z, 1), upos);  
}
```

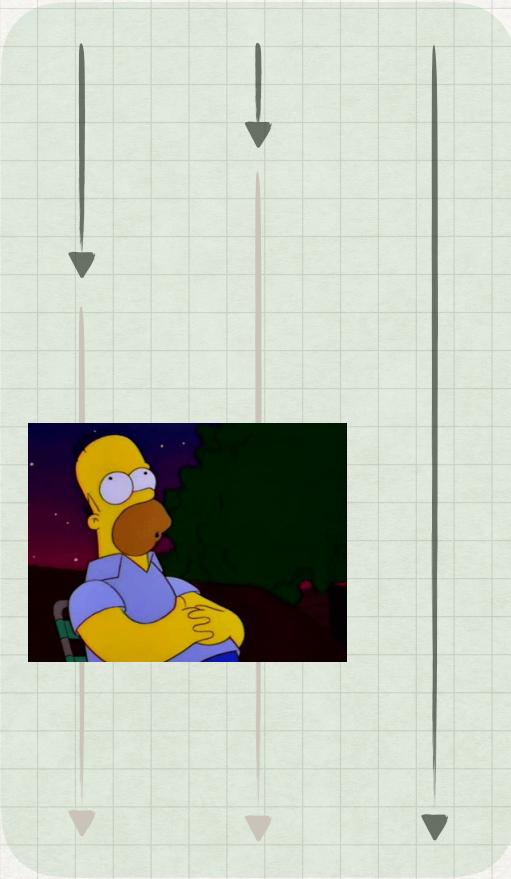
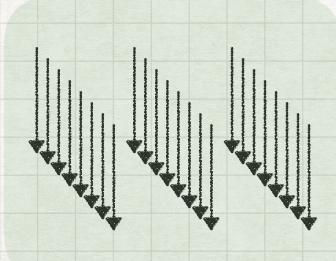
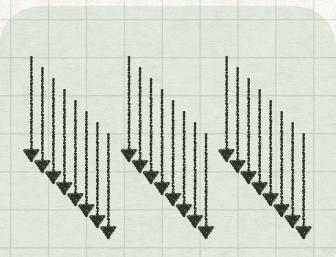
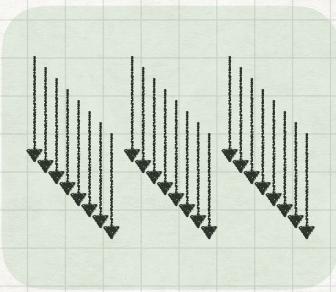
```
public func schlick(cosine: Float, indexOfRefraction: Float) -> Float {  
    var r0 = (1 - indexOfRefraction) / (1 + indexOfRefraction)  
    r0 *= r0  
    return r0 + ((1 - r0) * powf(1 - cosine, 5))  
}
```

```
static float schlick(float cosine, float indexOfRefraction) {  
    float r0 = (1 - indexOfRefraction) / (1 + indexOfRefraction);  
    r0 *= r0;  
    return r0 + ((1 - r0) * pow(1 - cosine, 5));  
}
```

```
struct Sphere {  
    var materialType: Int32 = 0  
    var color = Color()  
    var center = Vector()  
    var radius: Float = 0  
    var fuzz: Float = 0  
    var indexOfRefraction: Float = 1  
}
```

```
enum MaterialType {  
    lambertian, metalic, dialectric  
};  
  
struct Sphere {  
    MaterialType materialType = lambertian;  
    float3 color = float3(1);  
    float3 center = float3(0);  
    float radius = 0;  
    float fuzz = 0;  
    float indexOfRefraction = 1;  
};
```



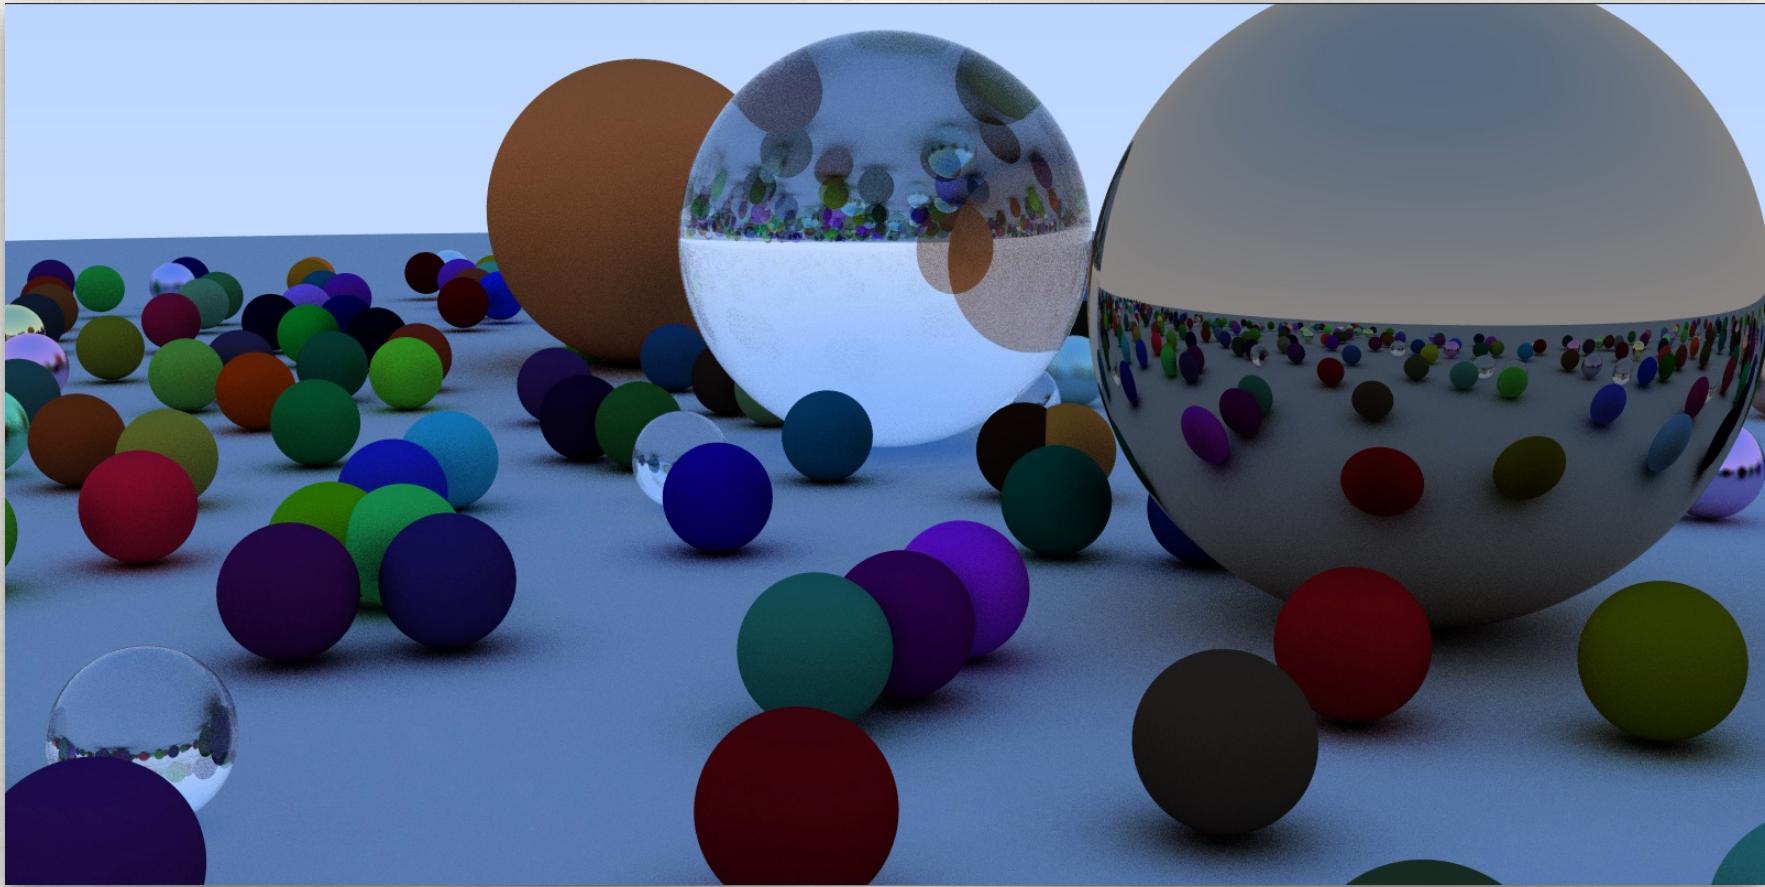


GOTCHAS & TIPS

- Artificial limit your loops on the metal side
- There are no random number generators native to metal
- Byte alignment will matter when passing structs back-n-forth
- There is no subclassing in MSL

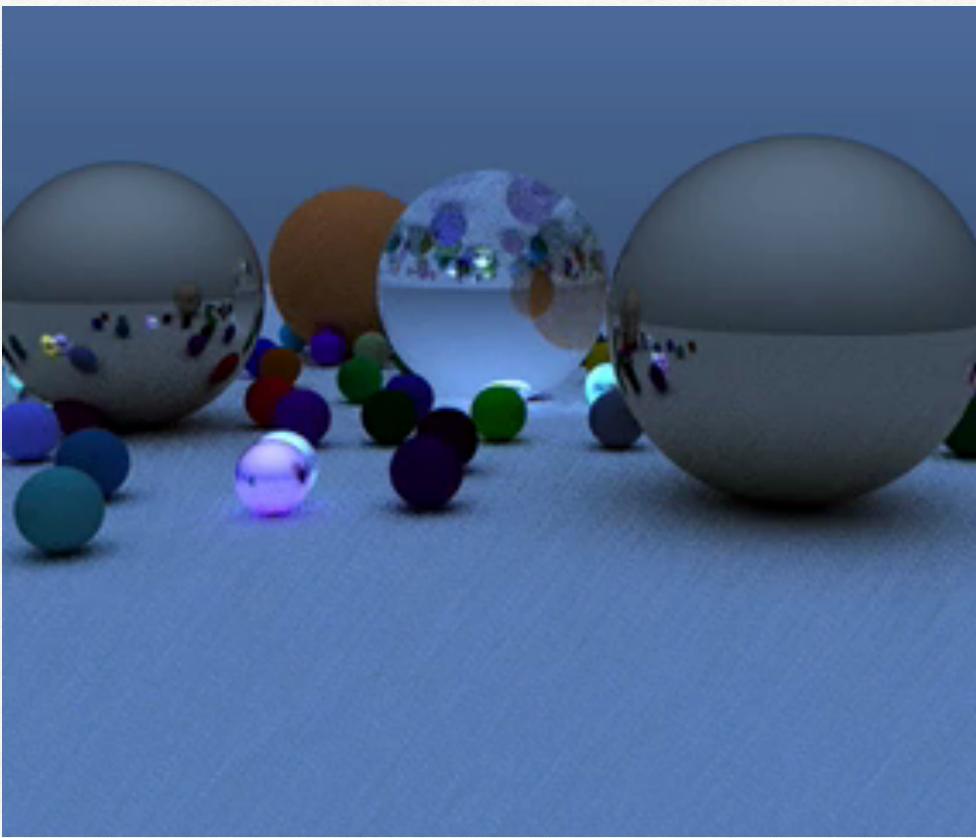
PROTOTYPING IN SWIFT

- Prototype in a real app so you can maximally debug it
- Avoid subclassing, or at least be aware that there is none in MSL
- If you need random numbers, pass them in or find an RNG to use
- Use **simd** framework



2012 MacBook Pro: ~60x

2013 MacPro: ~200x



Thanks!

Jonathan Blocksom
@jblocksom

Jeff Biggus
@hyperjeff



Code and slides up at:

<https://github.com/hyperjeff/360iDev-2016-RayTracing>