



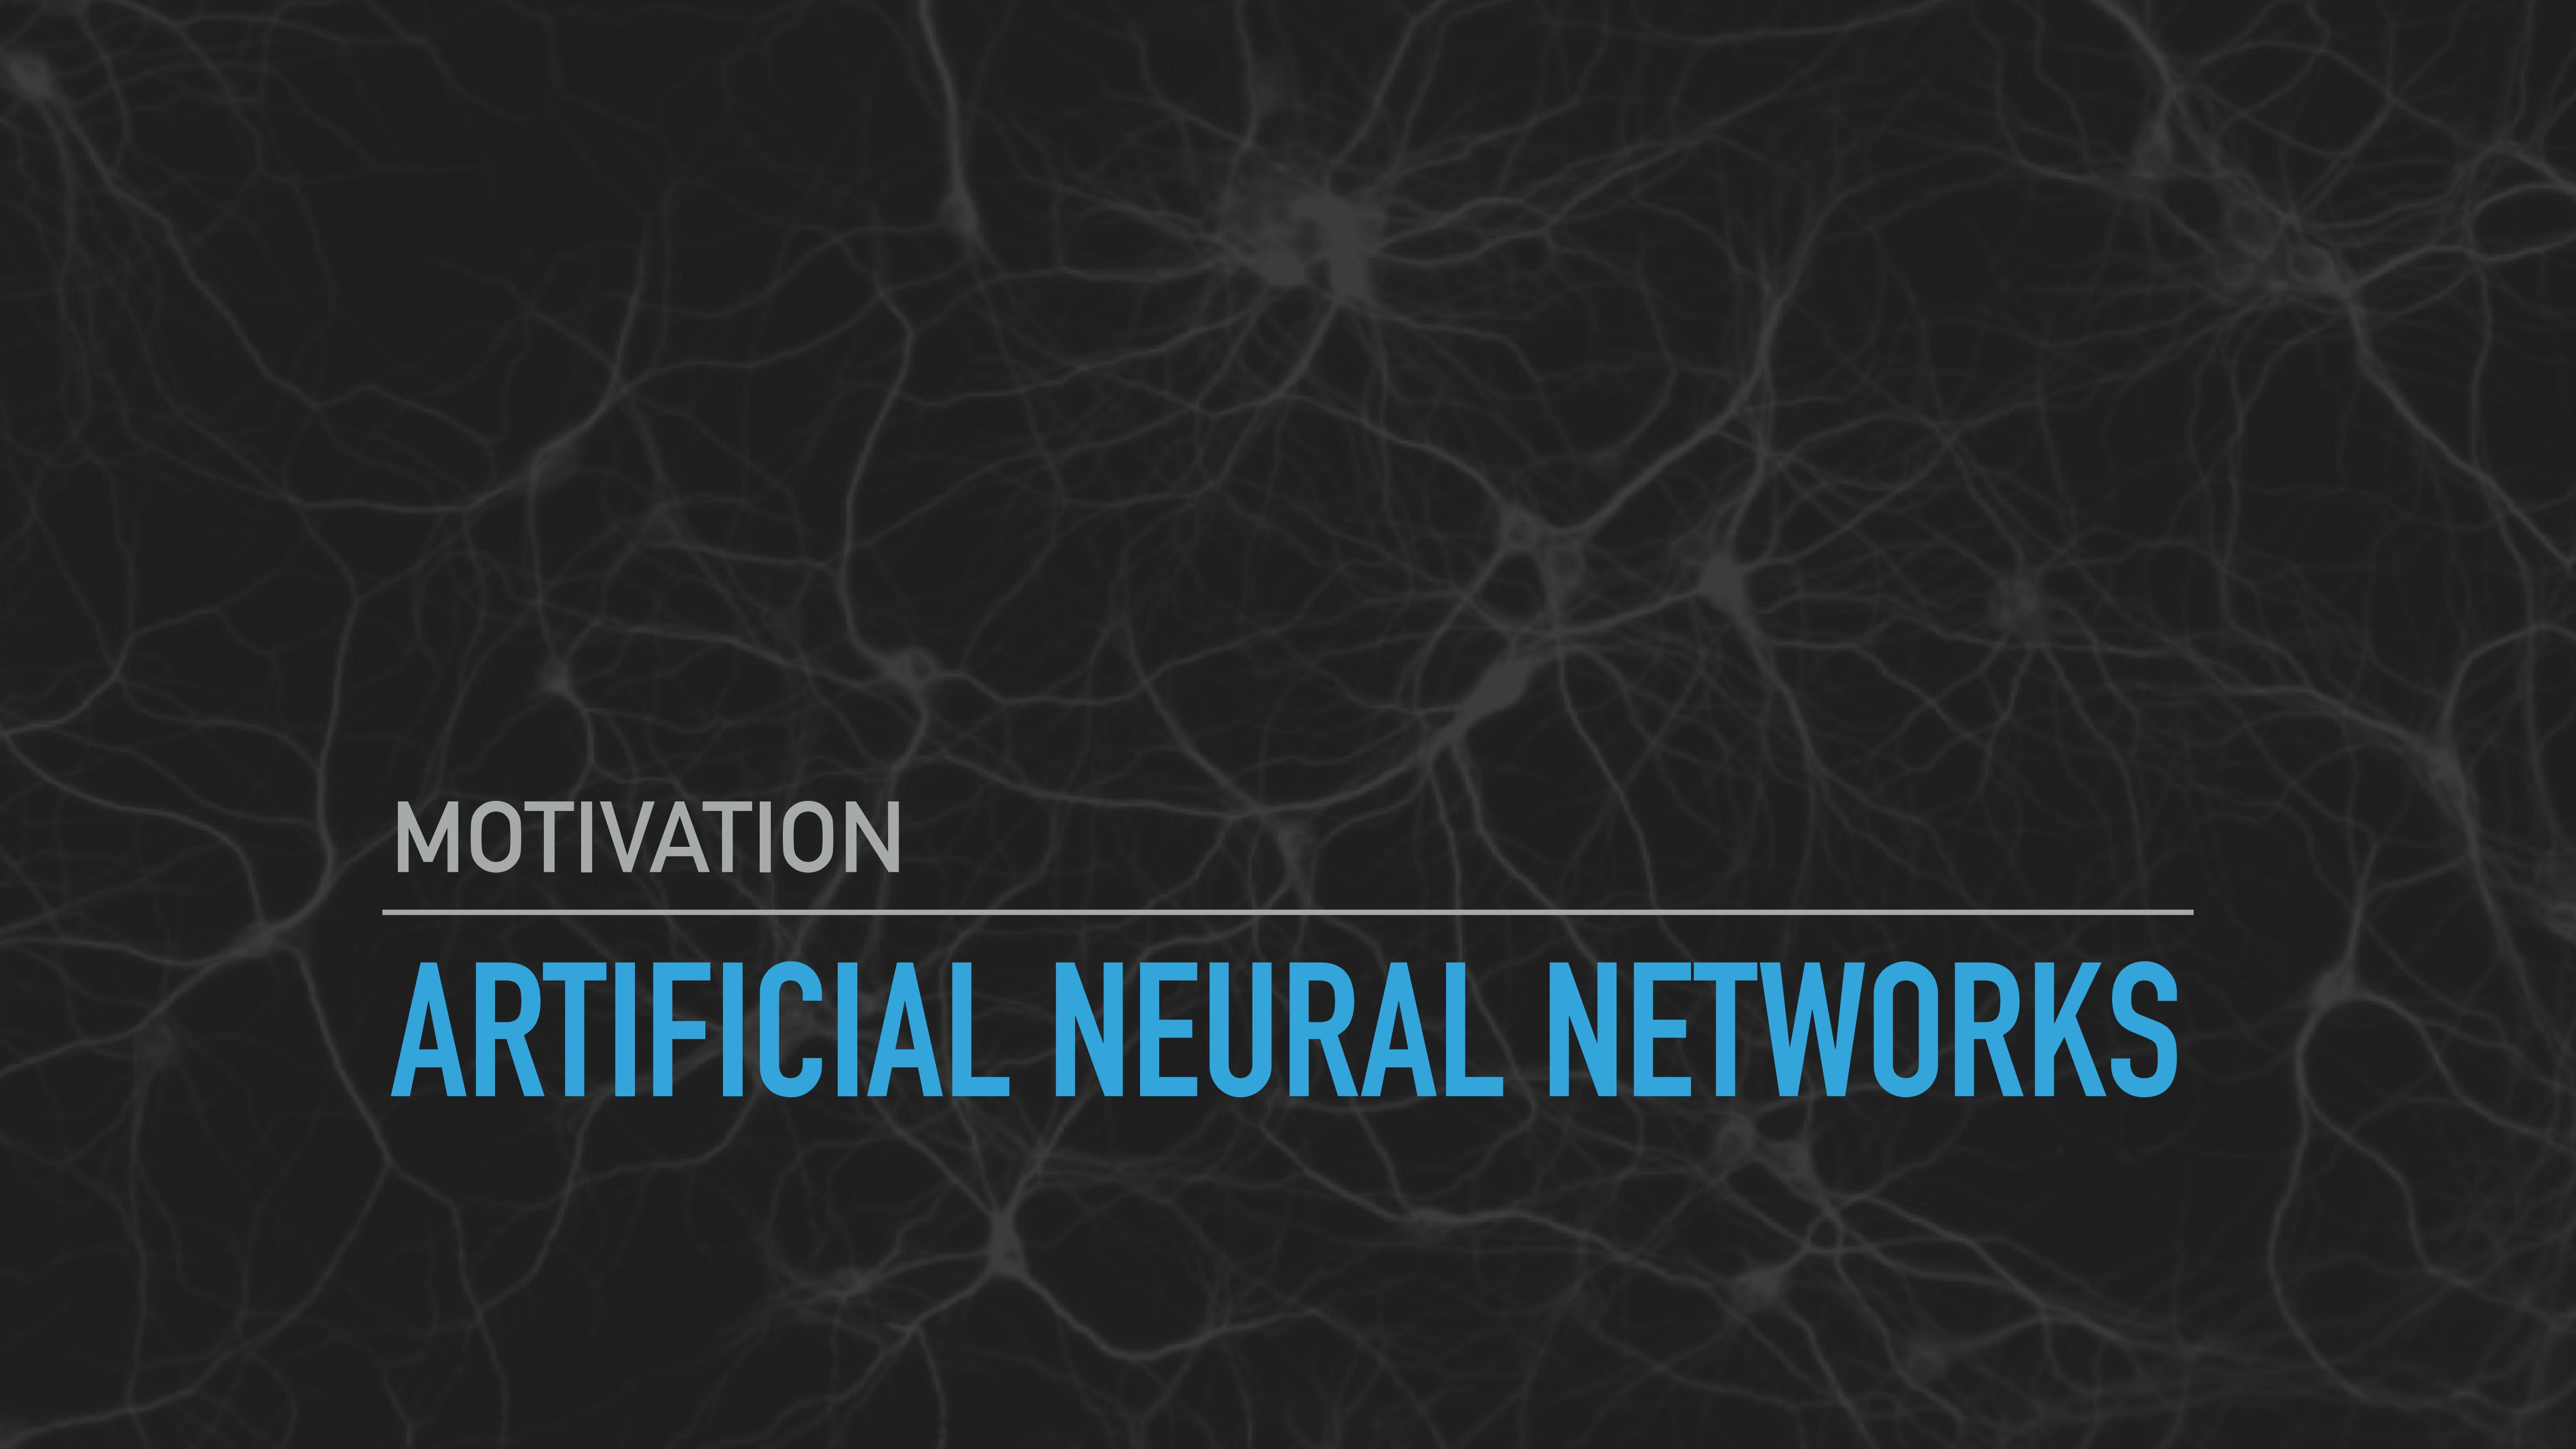
Chris Parrish
chris@twenty3.me
@twenty3

EXPLORING NEURAL NETWORKS

FOR SHAPE DETECTION

OVERVIEW

- ▶ Motivation
- ▶ Introduction to Neural Networks
- ▶ Building a Prototype
- ▶ What's next?



MOTIVATION

ARTIFICIAL NEURAL NETWORKS

EVOLUTION IN USER INTERFACE

- ▶ Macintosh popularized Point and Click
- ▶ iPhone brought tap, swipe and pinch
- ▶ What's next?

NEXT GENERATION INTERFACE

- ▶ Location and Positioning - GPS, iBeacon/Bluetooth LE, Magnetic Positioning
- ▶ Voice Assistants - Siri, Alexa, Cortona, Google Now
- ▶ Haptics, Handwriting, Drawing, Advanced Gestures

ARTIFICIAL INTELLIGENCE / MACHINE LEARNING

- ▶ Require some component of interpreting intent or meaning from the input signals
- ▶ Natural Language Understanding
- ▶ Computer Vision
- ▶ Prediction

CONSUMER HARDWARE ADVANCES

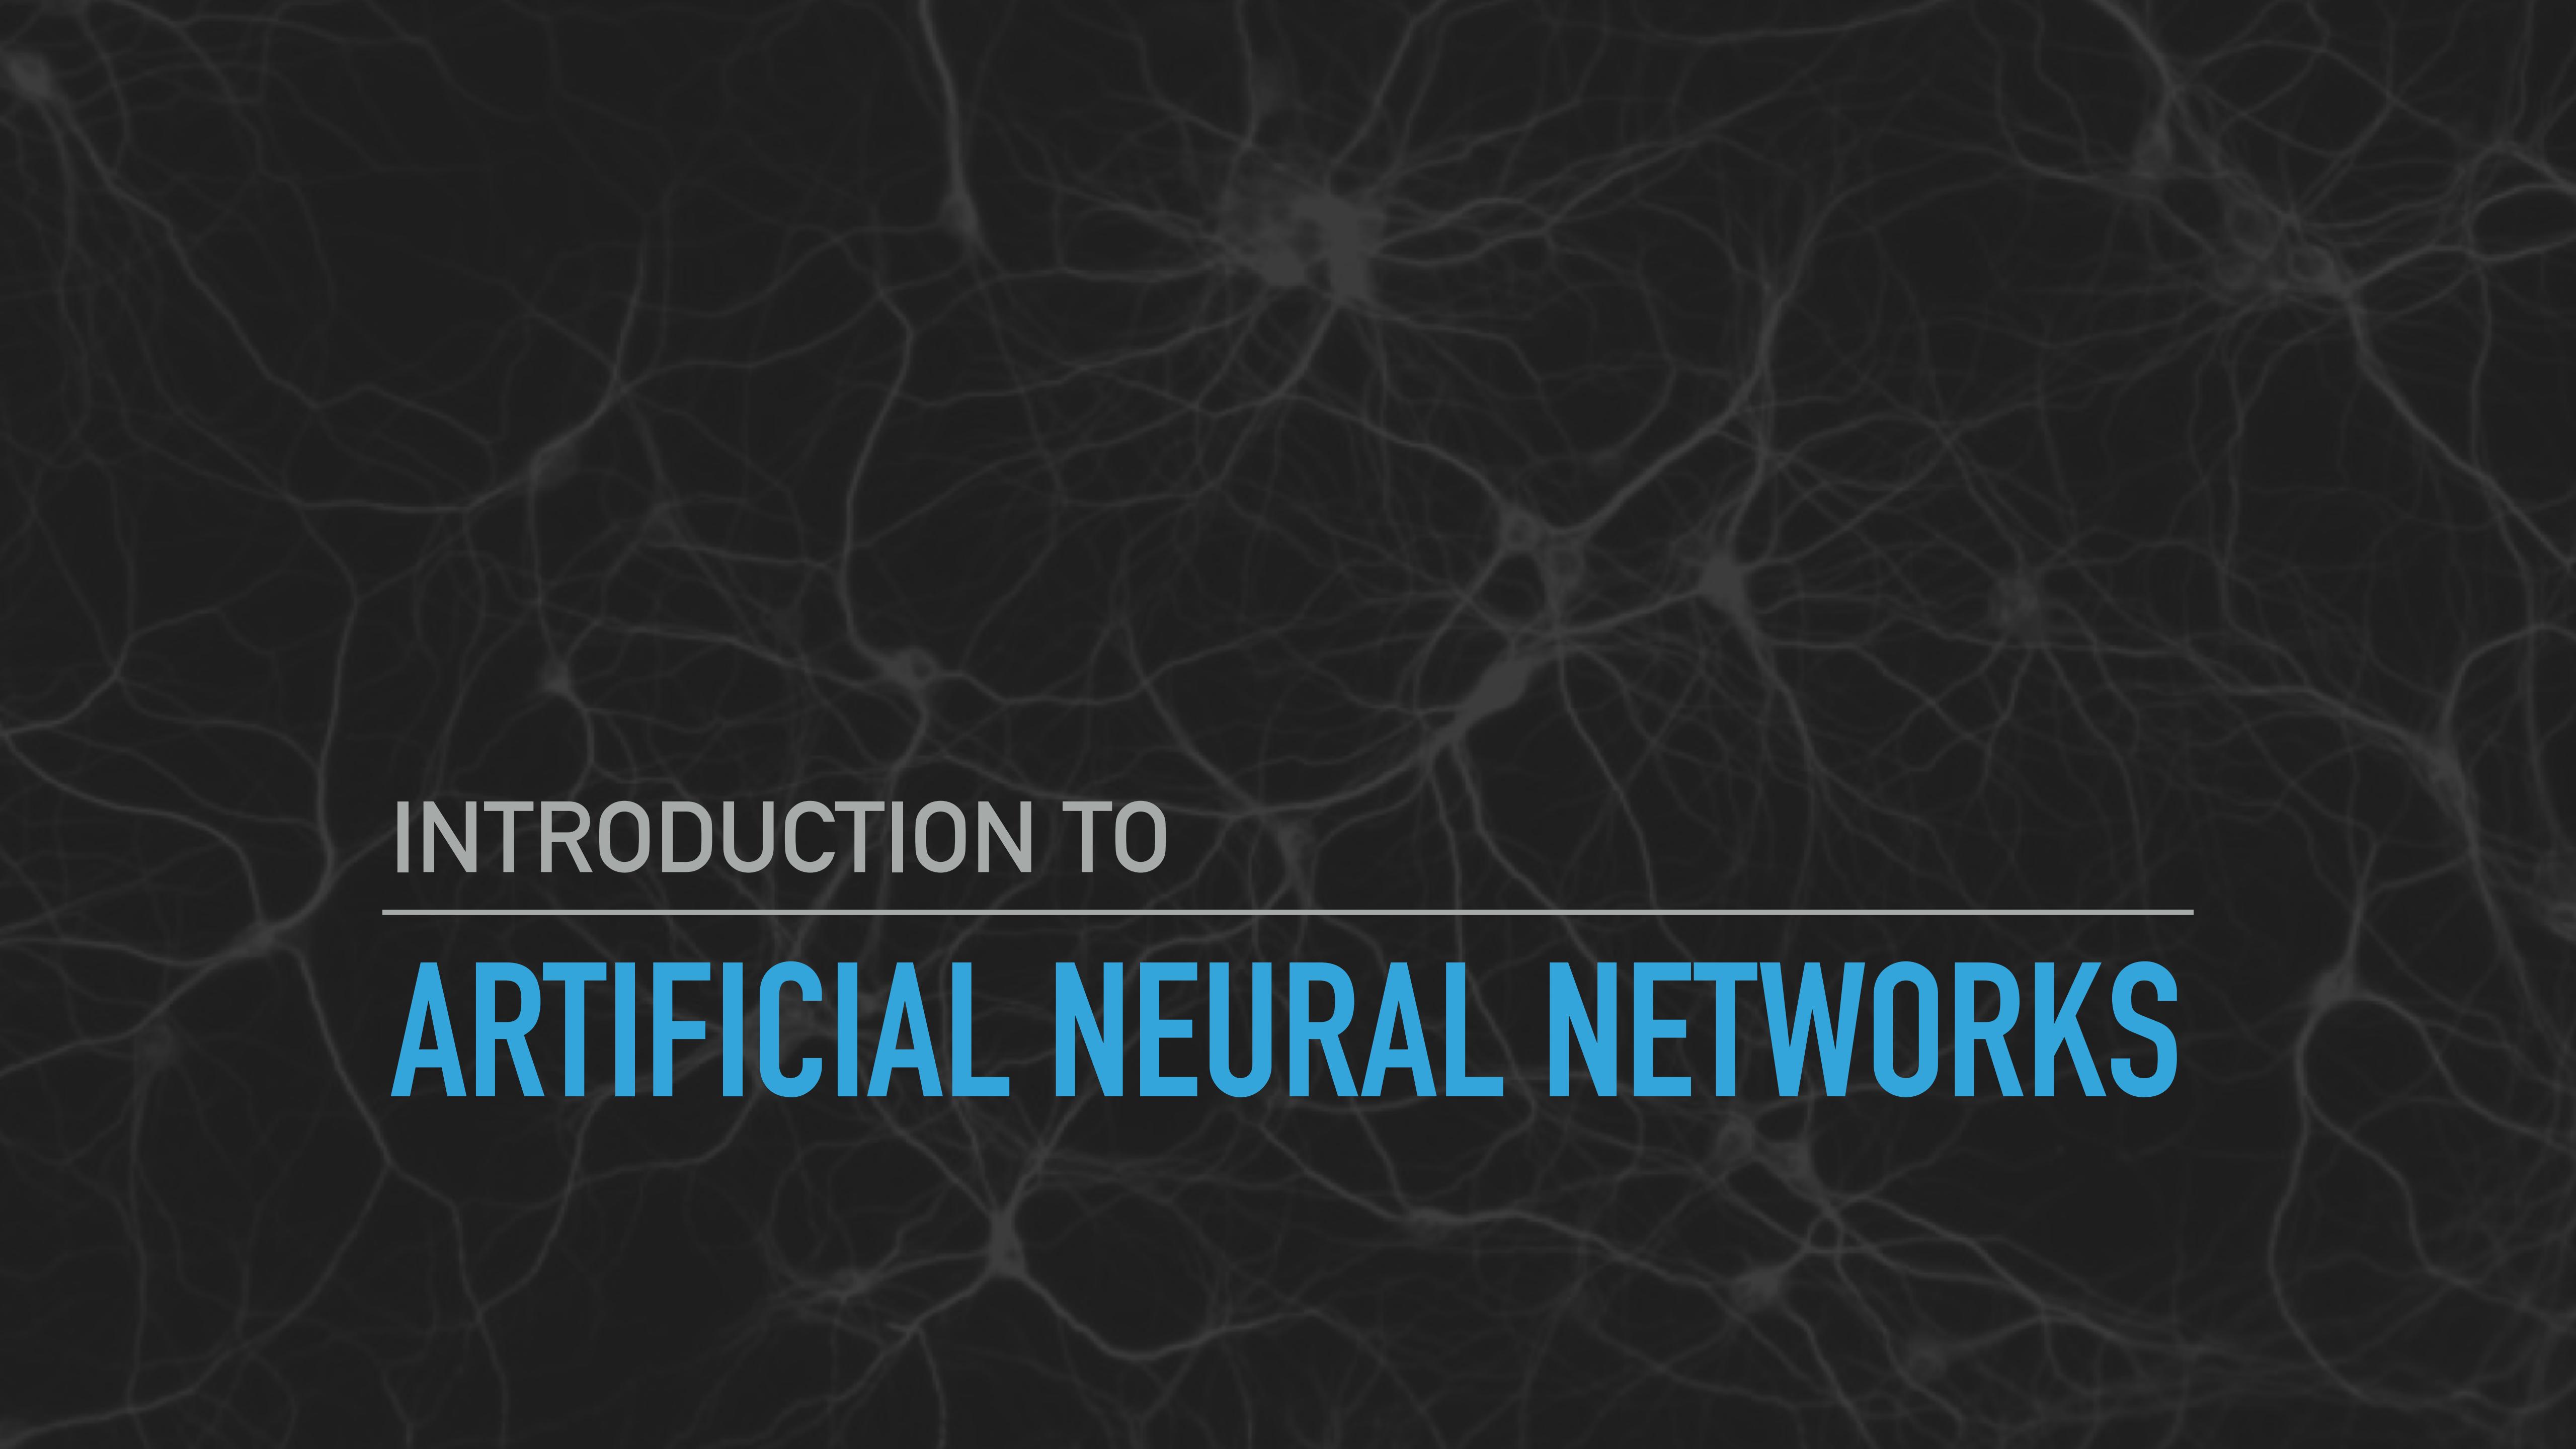
- ▶ Processing Power
- ▶ Mobile “workstations”
- ▶ Realistic to apply machine learning techniques in consumer applications
- ▶ New devices are ramping up the sensors and input methods
- ▶ iPad Pro
 - ▶ High resolution, low latency touch screen
 - ▶ Pencil

GESTURES

- ▶ Haven't advanced much beyond tap, swipe and pinch
- ▶ 'Hidden' and conflicting
- ▶ Drawing and handwriting are 'natural' skills
- ▶ Pencil gives us a great opportunity to explore a natural interface on iOS

NAPKIN DEMO





INTRODUCTION TO

ARTIFICIAL NEURAL NETWORKS

ARTIFICIAL NEURAL NETWORK

- ▶ Machine Learning technique
- ▶ Simple model inspired by biological nervous systems / brain
- ▶ Low-level compared to high-level, rule-driven Expert System
- ▶ Excellent at pattern recognition
- ▶ Adaptive / capable of learning

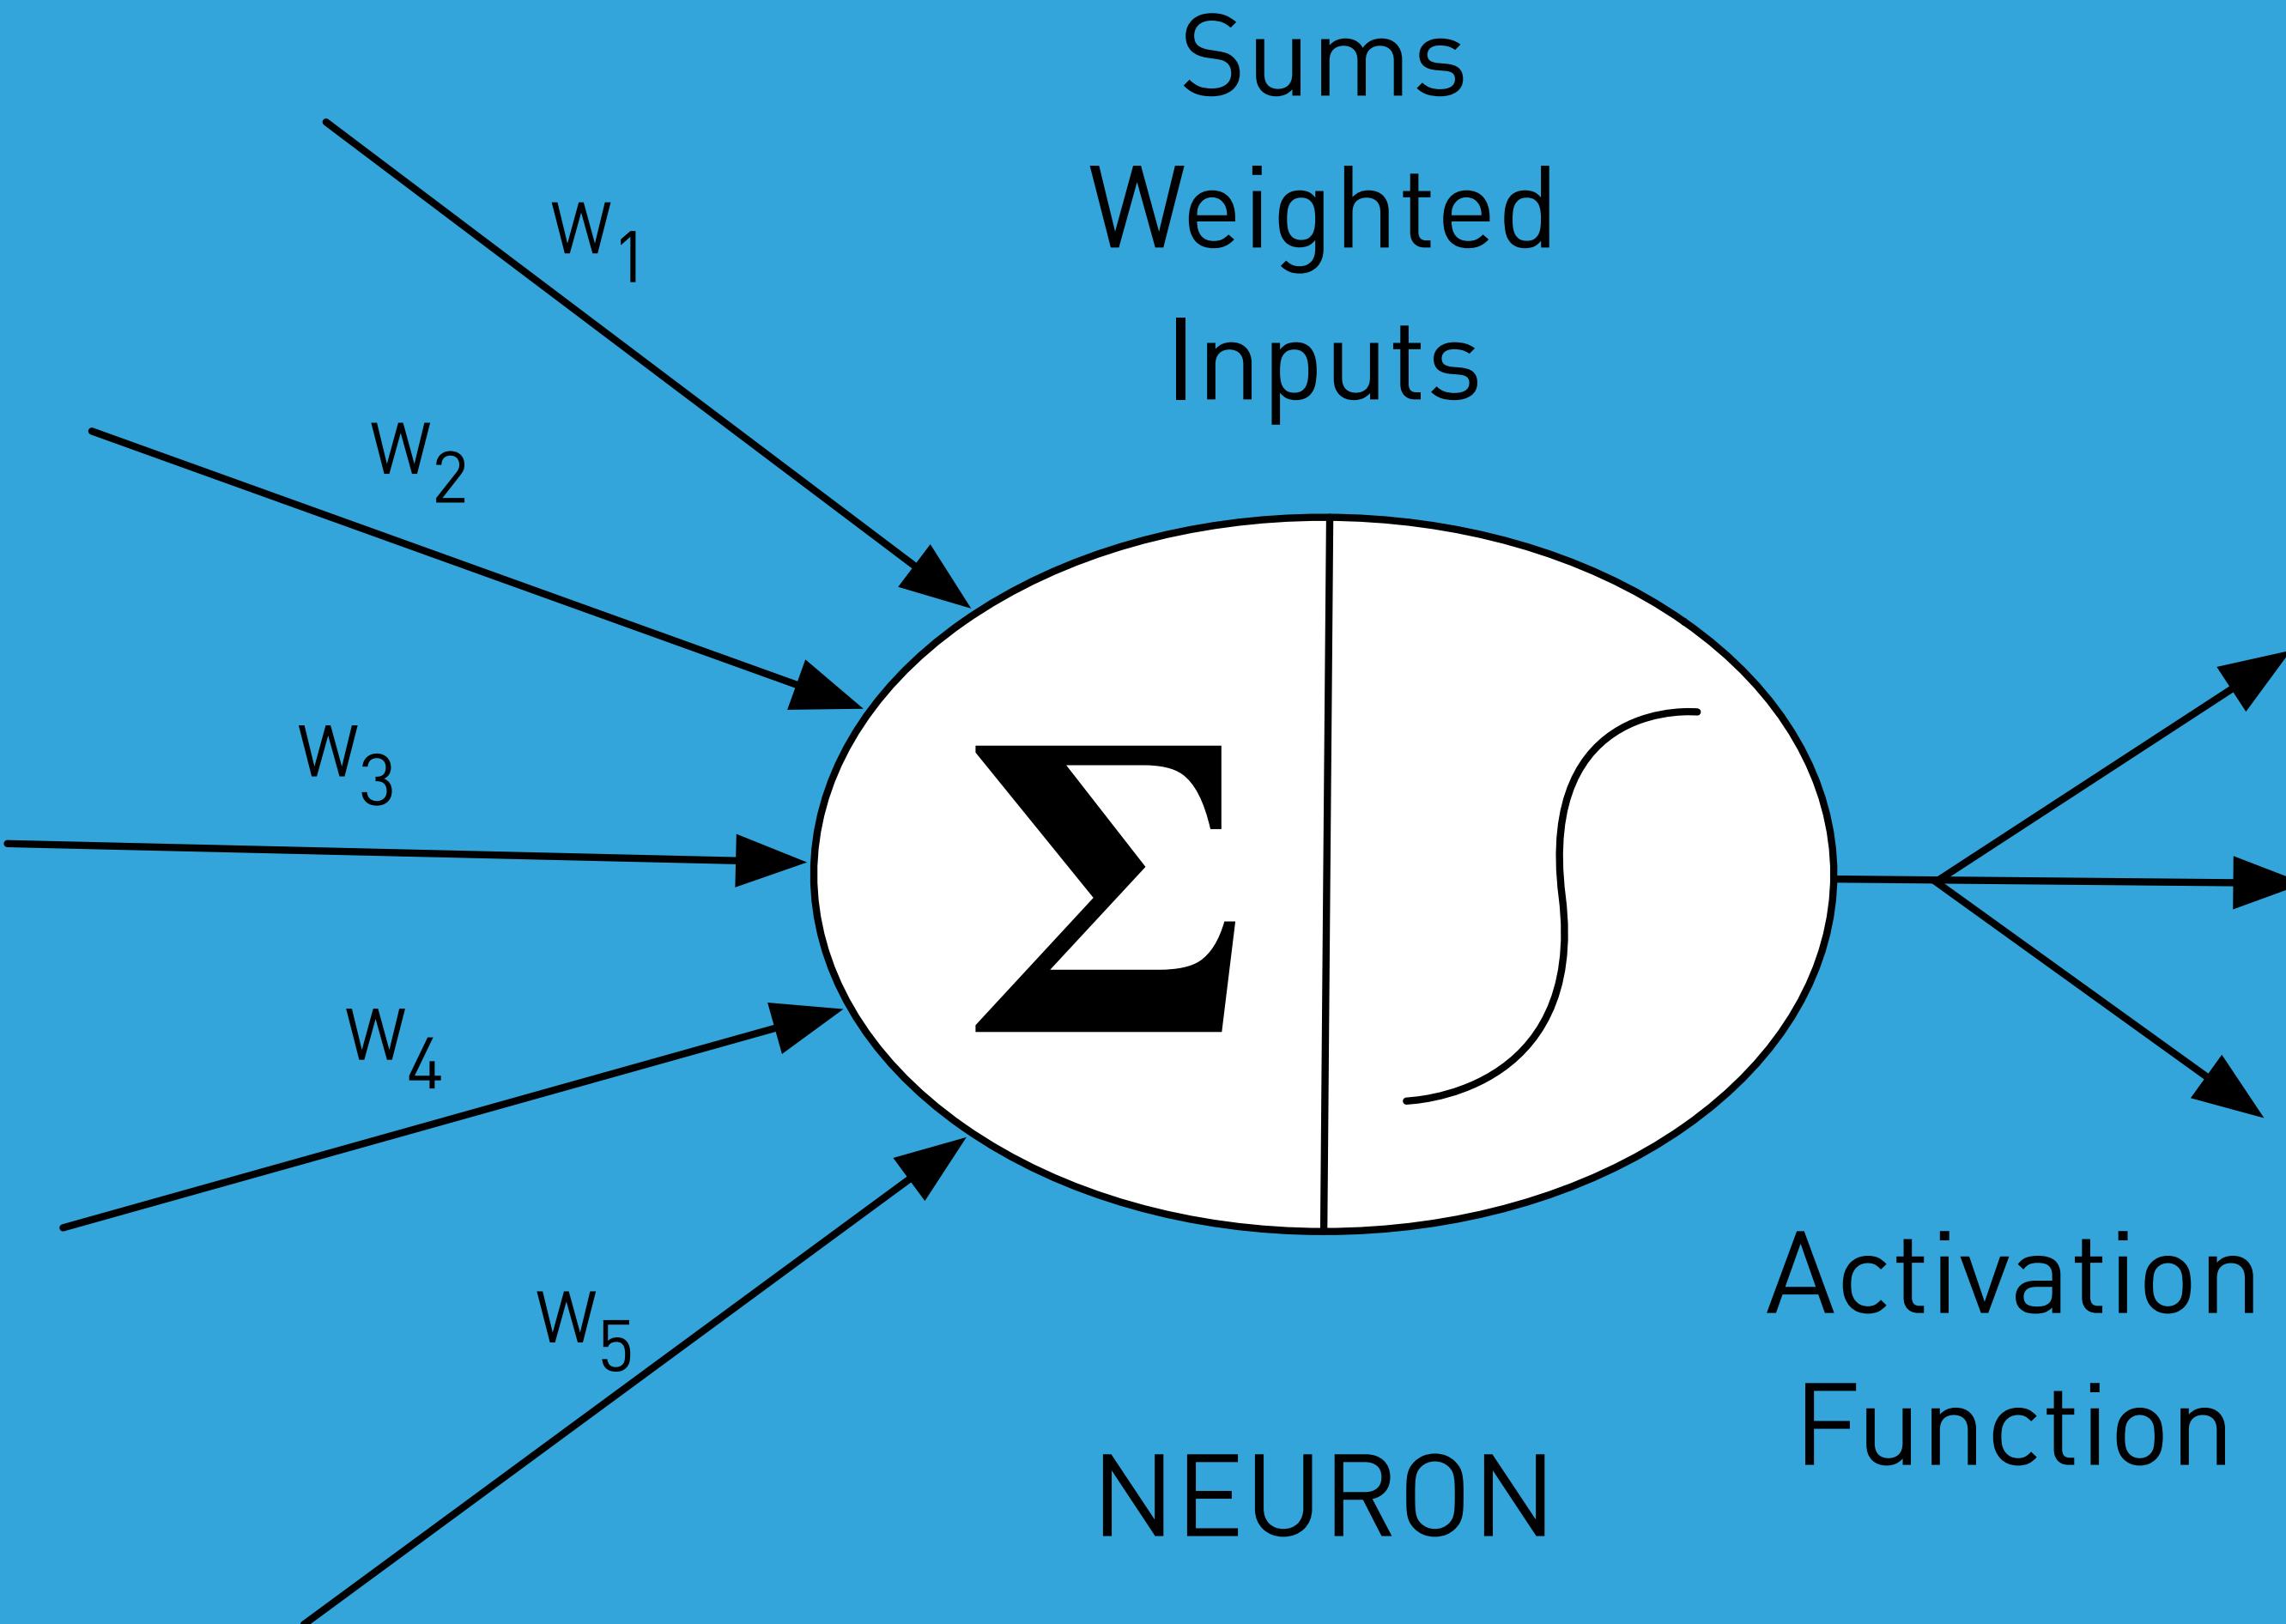
PROBLEMS WELL SUITED TO NEURAL NETWORKS

- ▶ Problems with a large number of inputs
- ▶ Problems that are hard to solve with rules-based systems
 - ▶ Image recognition
 - ▶ Handwriting recognition
 - ▶ Speech and language understanding
 - ▶ Autonomous driving / piloting

MODEL

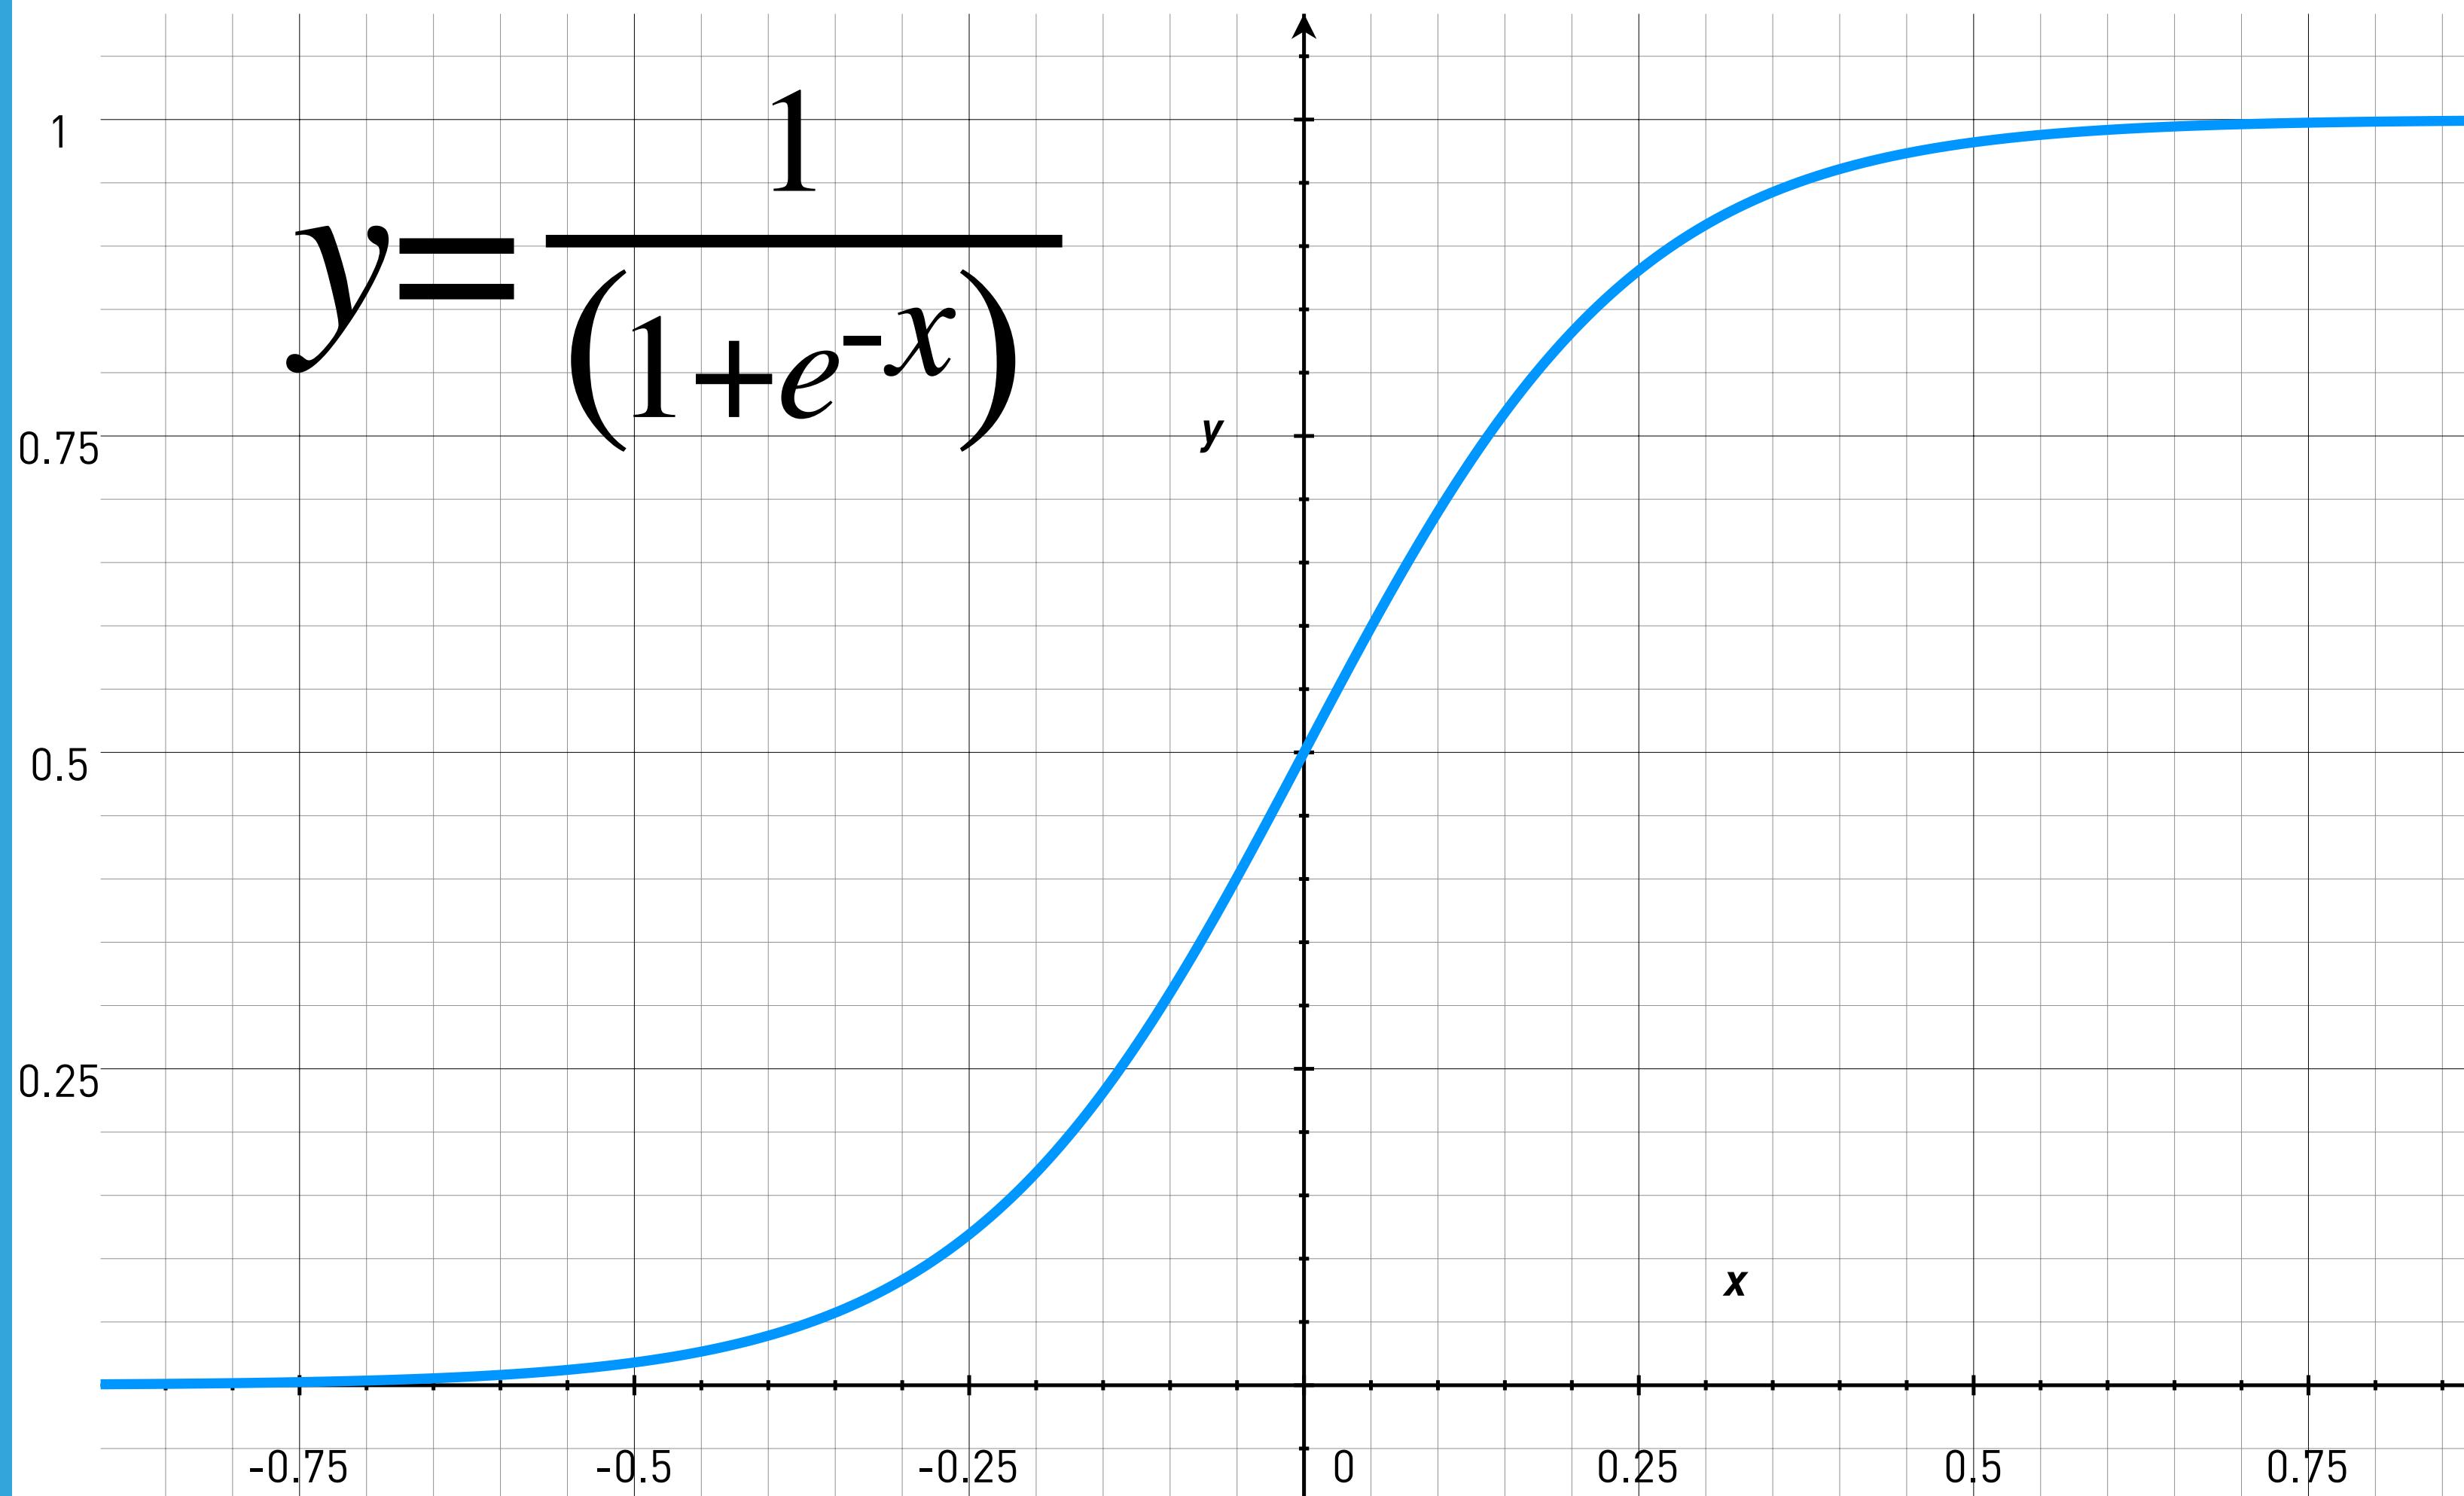
- ▶ NEURON
 - ▶ A node that sums incoming 'signals' and outputs a signal
 - ▶ Activation function controls output based on the inputs
- ▶ SYNAPSE
 - ▶ Connections between neurons
 - ▶ Synapse scales or 'weights' the input traveling along it

SYNAPSES



ACTIVATION FUNCTIONS

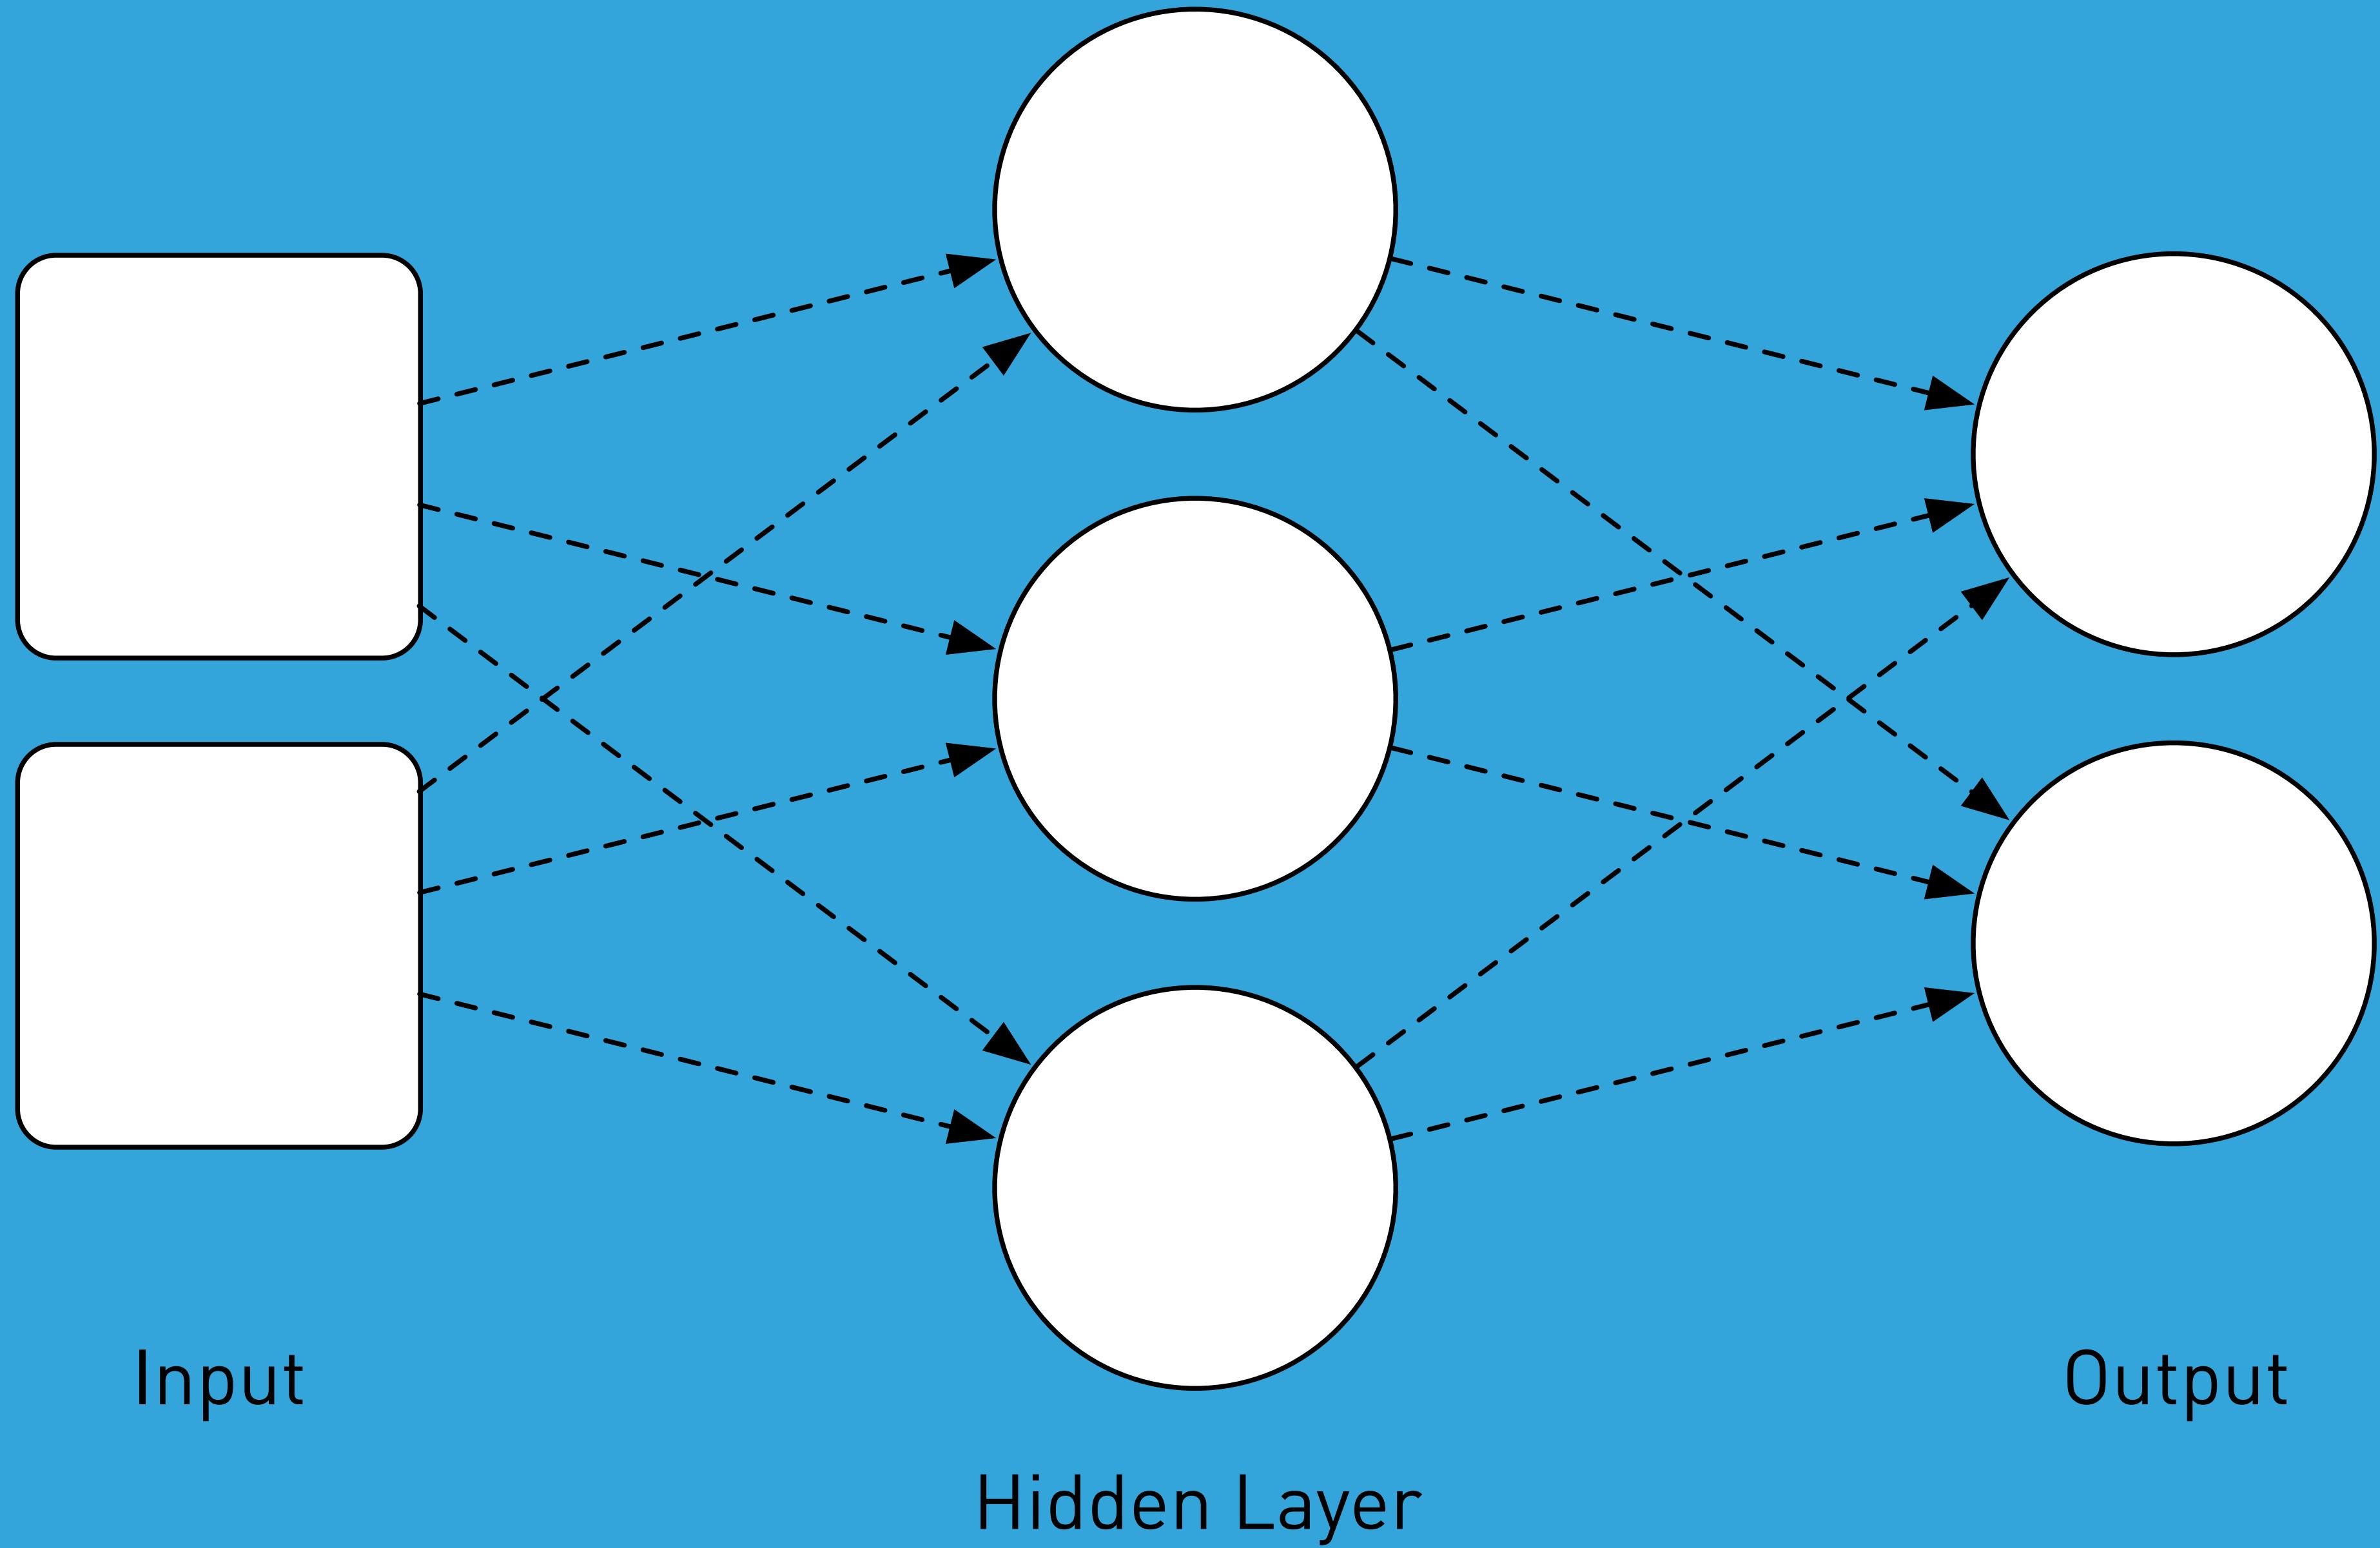
- ▶ A neuron's activation function controls how the inputs are converted into an output
- ▶ LINEAR
 - ▶ just scales the input
- ▶ THRESHOLD
 - ▶ If the sum of the inputs is greater than some X , activate
- ▶ SIGMOID
 - ▶ 'Squashes' the inputs to min and max (0, 1) or (-1, 1)

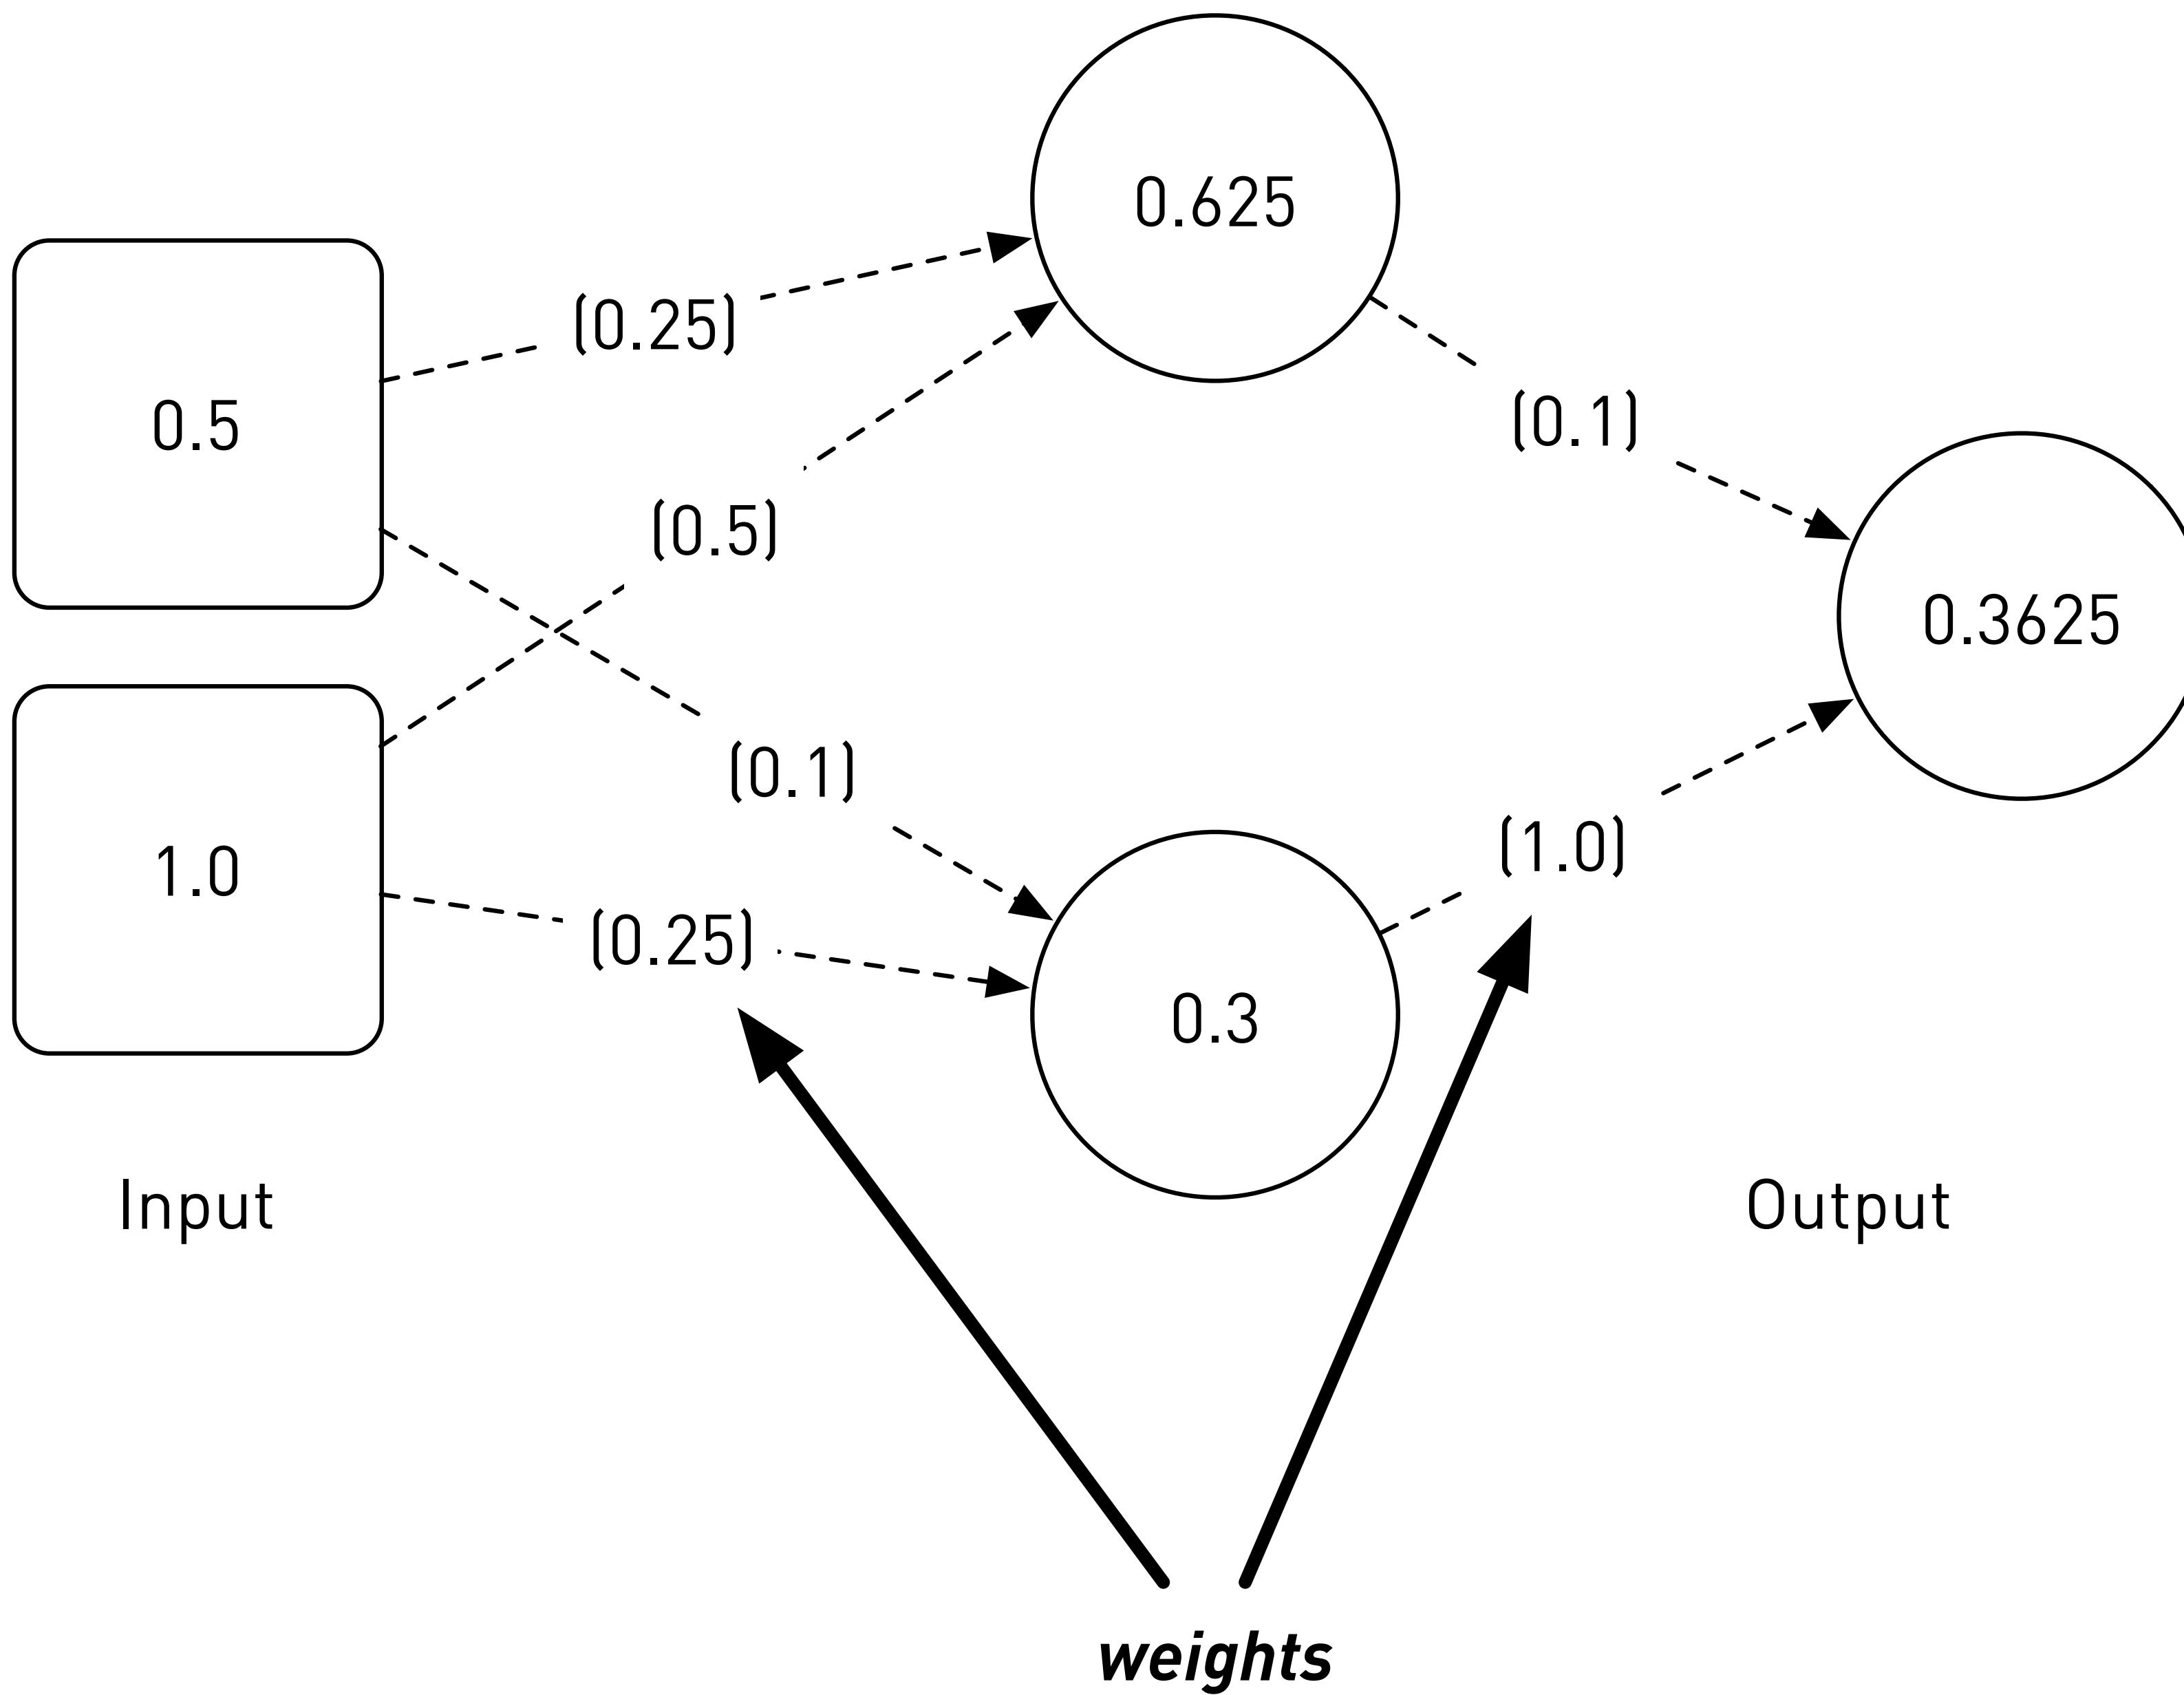


Sigmoid Activation Function

NETWORK

- ▶ Neurons connected in a network
 - ▶ Feed Forward network - signals only travel one direction (no feedback)
 - ▶ Layers between the input and output are called 'Hidden Layers'
 - ▶ Simple model: inputs, 1 hidden layer, outputs
 - ▶ Additional hidden layers allow approximating non-linear functions. Intuitively they classify higher levels of abstraction





OUTPUTS

- ▶ For classification problems (shapes, handwriting, etc) we have one output node per class (Circle, Square, Triangle)
- ▶ If an output node activates, we read that as the inputs being in that node's class
- ▶ Often, the output node's activation is between 0.0 and 1.0 and all nodes will have some value
- ▶ We can treat the node with the maximum activation as the answer to which class (confidence)

LEARNING

- ▶ The learning or adaptive nature of the network comes from adjusting the weights of the synapses between neurons.
- ▶ Network structure is unchanged
- ▶ In supervised learning we provide an input with a known answer and adjust the weights of each synapse to reduce the error
- ▶ Repeat with many examples from a training set to arrive at weights that can give correct answers for any input

ERROR (COST FUNCTION)

- ▶ How do we measure how wrong a wrong answer is?
- ▶ Square Error of some form generally

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (\hat{Y}_i - Y_i)^2$$

ADJUSTING WEIGHTS

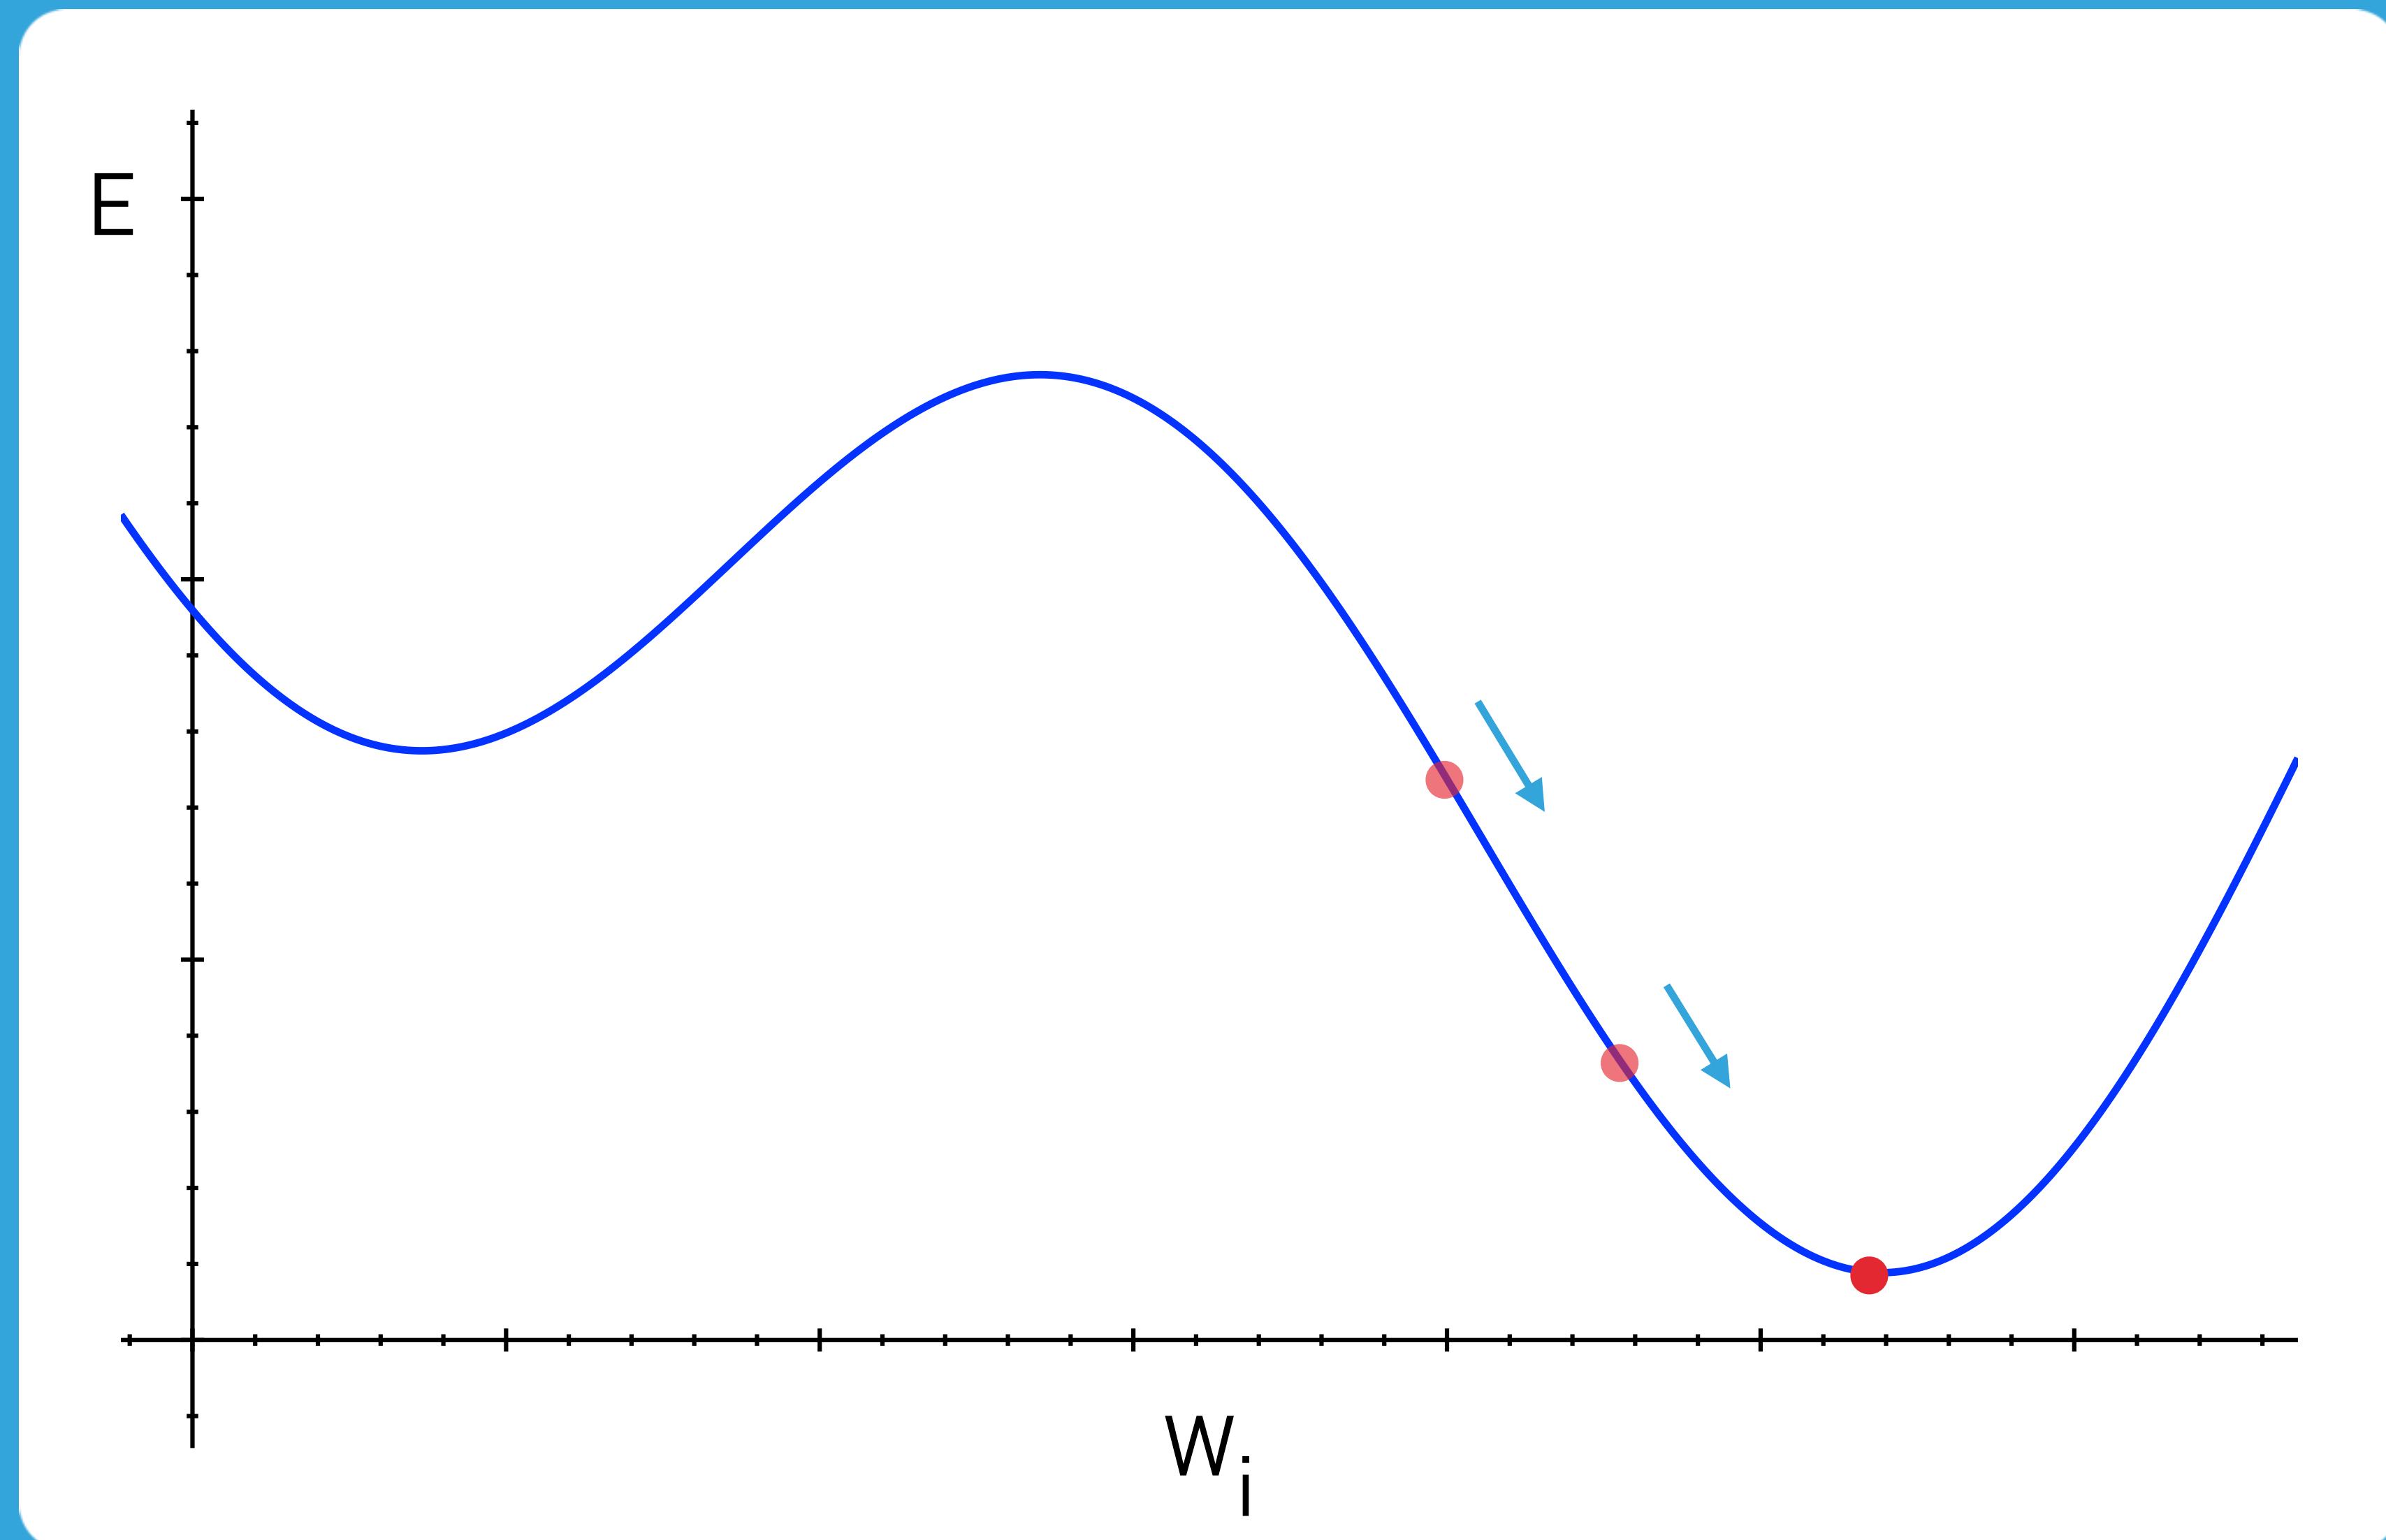
- ▶ How do we adjust the weights?
- ▶ How about just trying every possible weight and picking the ones that give the least error?
- ▶ For even the most simple networks, the computations required would require days, years, centuries, etc.

BACK PROPAGATION AND GRADIENT DESCENT

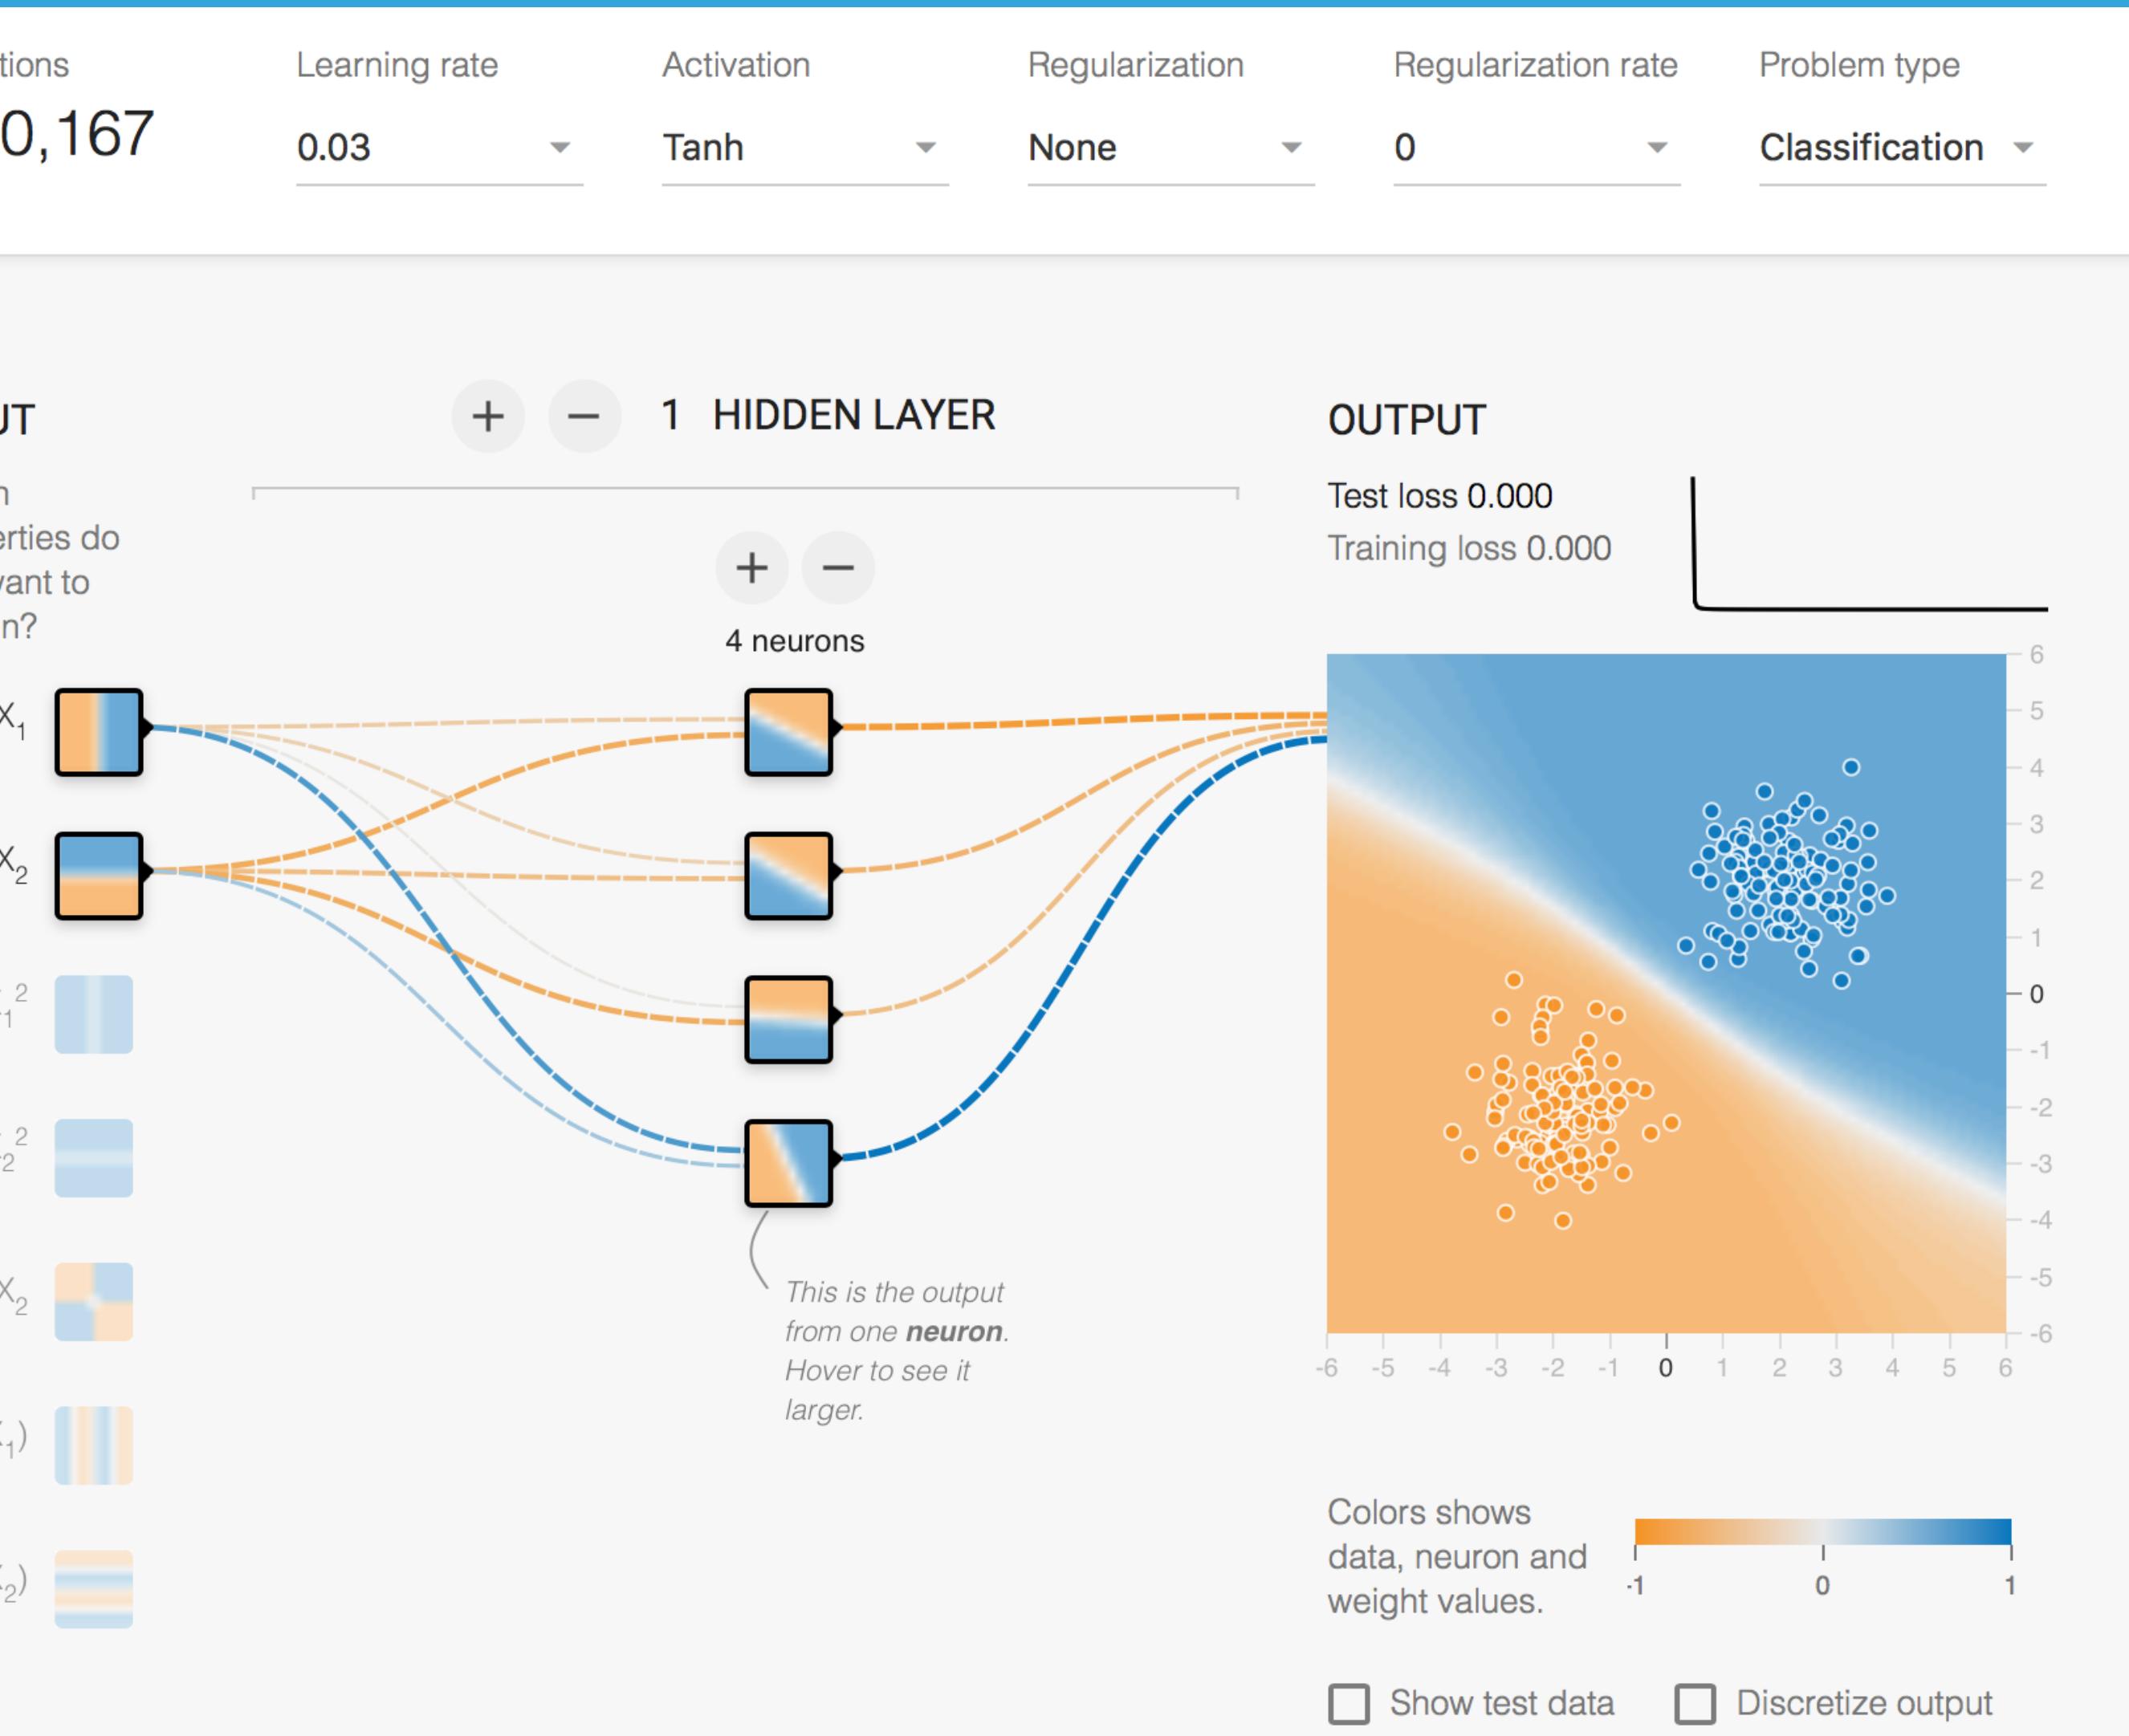
- ▶ Adjust the weights in small steps
- ▶ We want to go in the right 'direction' to make the error smaller. For each weight do we get larger or smaller?
- ▶ Construct the mathematical function of the error of the network in relation the inputs and weights
- ▶ Take the partial derivative (slope) of the error function with respect to each weight to determine which direction to adjust the weight

TUTORIALS

- ▶ [Gradient Descent](#)
- ▶ <https://www.youtube.com/watch?v=5u0jaA3qAGk>
- ▶ [Back Propagation](#)
- ▶ <https://www.youtube.com/watch?v=aVld8KMsdUU>



Error function for a weight



DEMO

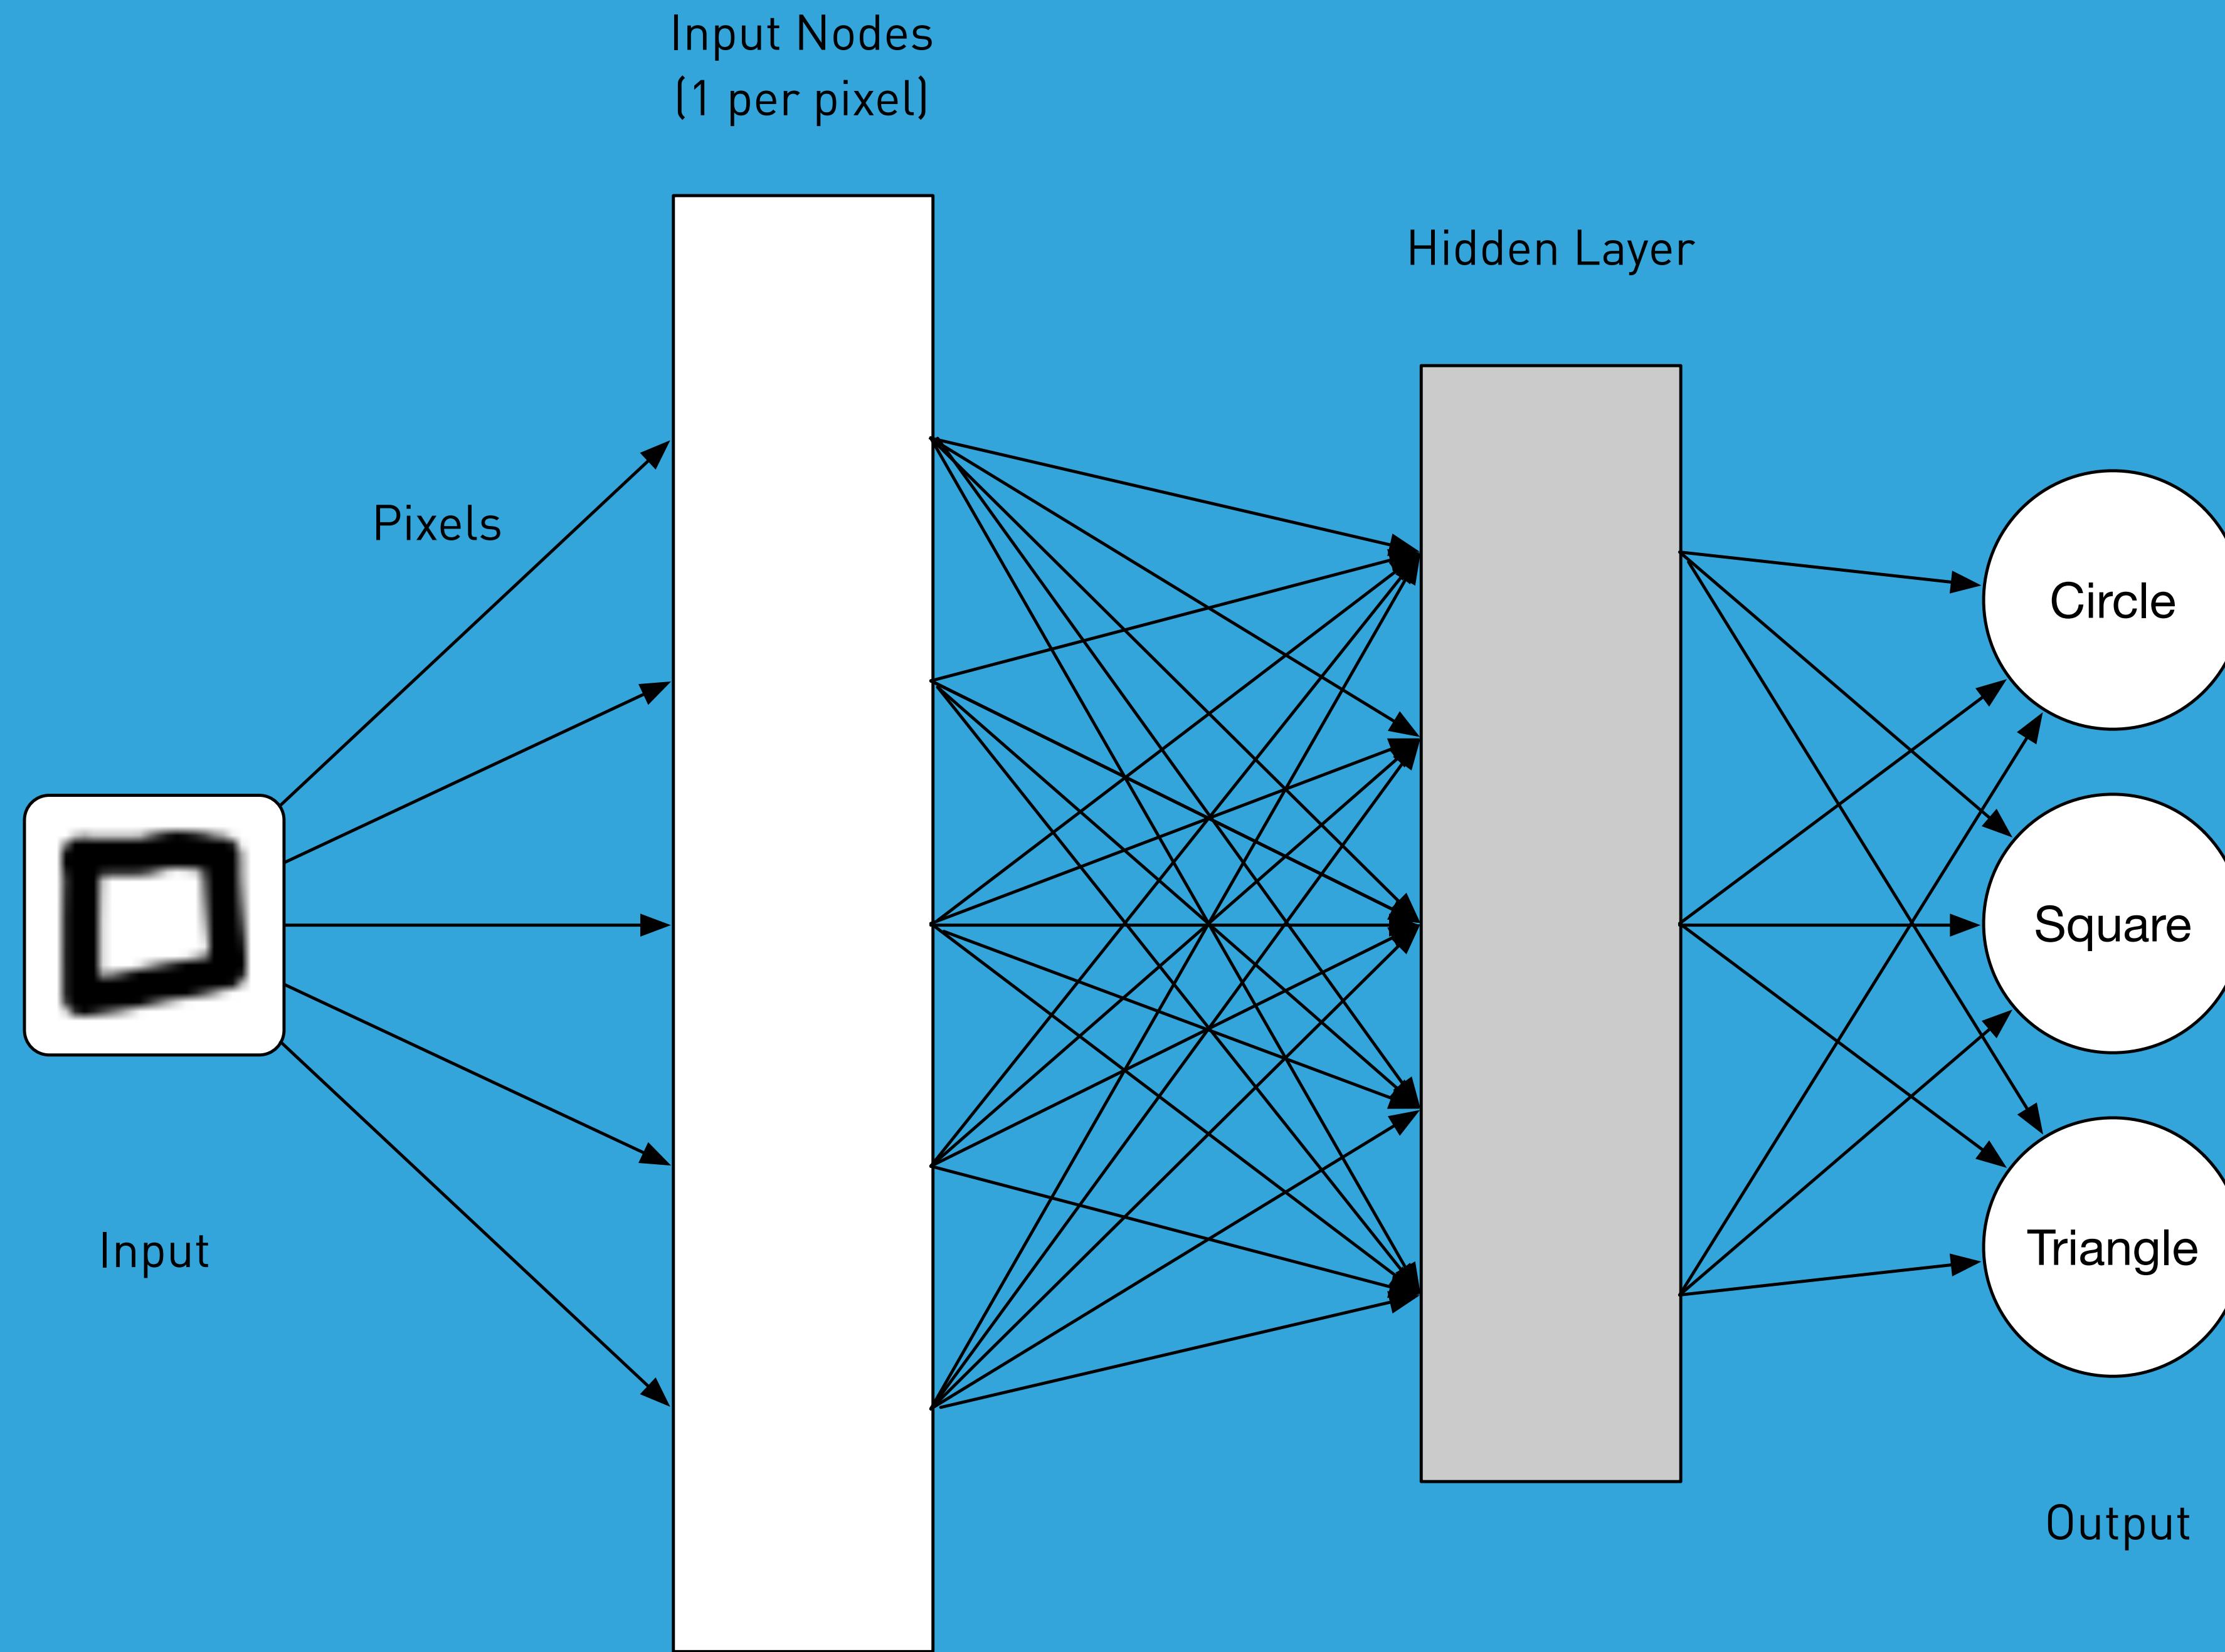
<http://playground.tensorflow.org>

PROTOTYPE

DETECTING DRAWN SHAPES

PROBLEM AND MODEL

- ▶ iOS iPad application
- ▶ Detect simple drawn shapes - Circle, Square, Triangle
- ▶ Feed shape 'data' into a Feed Forward Neural Network
- ▶ Output is confidence input matches a particular shape
- ▶ Supervised training on sample drawings



INPUT - MANY POSSIBLE SHAPE “FEATURES”

- ▶ Stroke count and direction
- ▶ Stroke velocity
- ▶ Number of angles
- ▶ Convexity
- ▶ **Image Bitmap**

OUTPUT

- ▶ 1 Output Node for each shape
- ▶ O_0 : Circle, O_1 : Square, O_2 : Triangle
- ▶ Resulting value 0.0 - 1.0 interpreted as confidence inputs
match that output node's shape
- ▶ If output == [0.0, 1.0, 0.0], 100% confidence input is a square

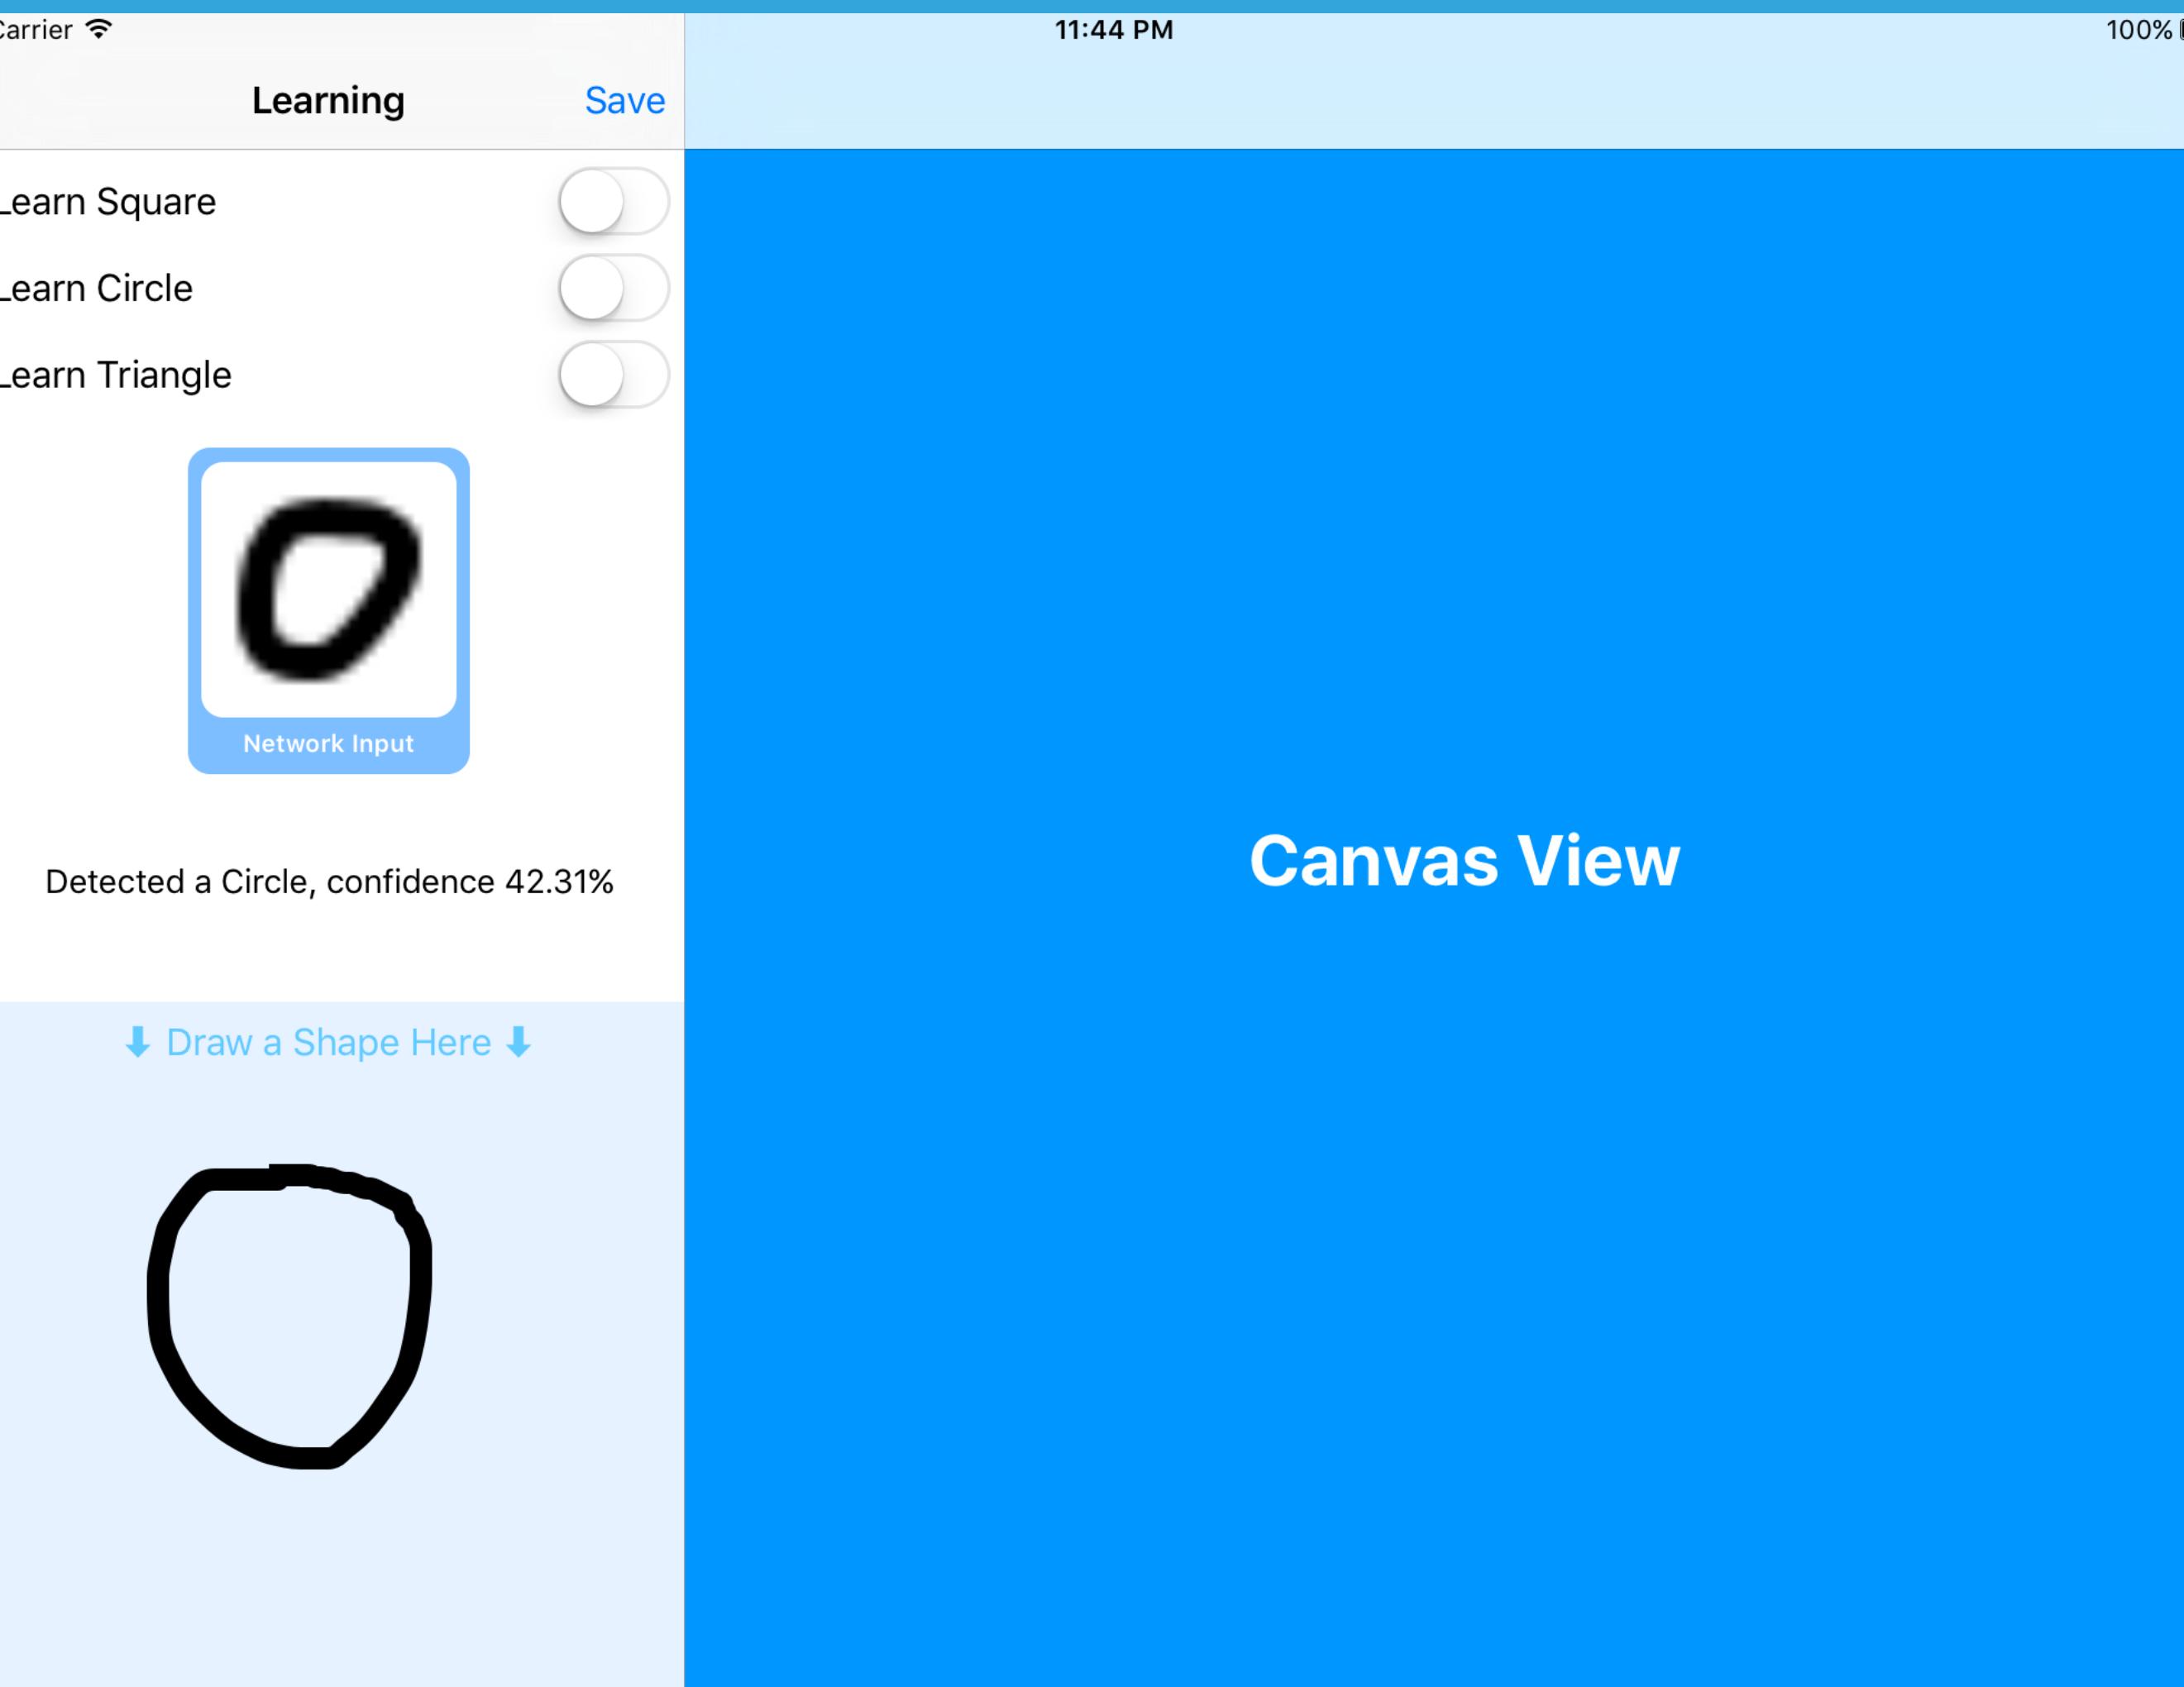
NETWORK STRUCTURE

- ▶ 3 Layer, full-connected network
- ▶ 1 input for each image pixel
- ▶ 1 output for each type of shape to classify

PARAMETERS

- ▶ 784 input nodes
 - ▶ 28 × 28 pixel bitmap of drawn shape
- ▶ 280 hidden nodes
- ▶ 3 Output nodes
- ▶ Sigmoid Activation Function
- ▶ Average Error Function

**TRIAL, ERROR,
GUESSING –
SCIENCE!**



DEMO

WE NEED SOME NEURONS AND SYNAPSES!

- ▶ The model for synapses and neurons isn't too complex,
We got this, couple of hours, easy!
- ▶ What about that activation function - Was his name Sigmund
or something?
Hmmm, we can probably find an example

WE NEED SOME NEURONS AND SYNAPSES!

- ▶ What about computing the error or cost?

Lets just use that squared error, not to bad right?

- ▶ We also need Gradient Descent to adjust the weights- partial derivatives, matrix math, dot products, summation, local minima, non-convex functions

OKAY... THAT'S A LOT OF MATH.
MAYBE MY ESTIMATE WAS OFF A LITTLE BIT.

SWIFT AI

- ▶ Fortunately there are a lot of open source libraries that provide basic Neural Network implementations
- ▶ Swift-AI to the rescue
- ▶ <https://github.com/collinhundley/Swift-AI>
- ▶ Compact and simple interface

OVERVIEW OF SWIFT AI INTERFACE

```
class FFNN {  
    public init(inputs: Int, hidden: Int, outputs: Int,  
                learningRate: Float = 0.7, momentum: Float = 0.4,  
                weights: [Float]? = nil,  
                activationFunction: ActivationFunction = .Default,  
                errorFunction: ErrorFunction = .Default(average: false))  
}
```

OVERVIEW OF SWIFT AI INTERFACE

```
// Propagates the given inputs through the neural network,  
// returning the network's output.  
// - Parameter inputs: An array of `Float`s, each element  
// corresponding to one input node.  
// - Returns: The network's output after applying the given  
// inputs, as an array of `Float`s.  
public func update(inputs: [Float]) throws -> [Float]
```

OVERVIEW OF SWIFT AI INTERFACE

```
    /// Trains the network by comparing its most recent output to the  
given 'answers', adjusting the network's weights as needed.  
    /// - Parameter answer: The 'correct' desired output for the most  
recent update to the network, as an array of `Float`s.  
    /// - Returns: The total calculated error from the most recent  
update.
```

```
public func backpropagate(answer answer: [Float]) throws -> Float
```

OVERVIEW OF SWIFT AI INTERFACE

```
/// Reads a FFNN from file.  
/// - Parameter filename: The name of the file, located in the  
default Documents directory.  
public static func fromFile(filename: String) -> FFNN?  
  
/// Writes the FFNN to file.  
/// - Parameter filename: The name of the file to write to. This file  
will be written to the default Documents directory.  
public func writeToFile(filename: String)
```

CAPTURING THE DRAWING

- ▶ UIPanGestureRecognizer
- ▶ Connect the touch points in a quadratic UIBezierPath
- ▶ Stroke the path
- ▶ Update and redraw path as new touches occur
- ▶ If no new touches after 1 second, consider the shape complete

WHY A QUADRATIC BEZIER PATH?

- ▶ We could just draw a 'brush tip' along the touch points instead of a bezier path
- ▶ Path gives us an easy method to smooth out touches to get a more natural result
- ▶ Path gives us a way to normalize the stroke width when pre-processing for the network (more later)
- ▶ Path doesn't lend itself to variable width strokes and produces odd effects on direction reversal

SMOOTHING THE PATH

- ▶ Use the mid-point between current touch and previous touch as the next path point
- ▶ Use the current and previous touch point as the bezier path control points
- ▶ See WWDC 2012 session - [Building Advanced Gestures Recognizers](#)

PREPARING THE BITMAP FOR PROCESSING

- ▶ Capture drawing as an image
- ▶ Crop to just the path bounds
- ▶ Normalizes the aspect (*rectangles become squares*)
- ▶ Scale the image down
 - ▶ just enough pixels to capture the essential features
 - ▶ Re-stroke the path so stroke-width is invariant

PREPARING THE BITMAP FOR PROCESSING

- ▶ Convert to grayscale
- ▶ RGB would mean we'd need 3 input nodes for each pixel
- ▶ Color is not important for our problem
- ▶ Grayscale vs. Monotone
 - ▶ Grayscale values influence the networks behavior
 - ▶ Network expect inputs between (0.0, 1.0)



XCODE TIME

HOW DO WE TRAIN THIS THING?

- ▶ Depending on the problem space and the network configuration, you might require hundreds or thousands of samples
- ▶ Feed samples with answers to network and let it go
- ▶ Any online corpuses for hand-drawn shapes?
- ▶ Drawing is fun!





CONCLUSION

WHERE DO WE GO NOW?

NEXT STEPS

- ▶ Rejection
- ▶ Skew, rotation
- ▶ Capture additional features
- ▶ Adaptive learning / user correction
- ▶ Performance

SOME RESOURCES THAT MIGHT BE USEFUL

- ▶ BrainCore <https://github.com/aleph7/BrainCore>
- ▶ SwiftSimpleNueralNetowrk <https://github.com/davecom/SwiftSimpleNeuralNetwork>
- ▶ AIToolBox <https://github.com/KevinCoble/AIToolbox>
- ▶ BirdBrain <https://github.com/jordenhill/Birdbrain>
- ▶ MLKit. <https://github.com/Somnibyte/MLKit>

The End